

**OUTER SPACE AND FOURIER SPACE:
UNDERSTANDING FOREGROUNDS FOR HI EPOCH OF REIONIZATION**

MEASUREMENTS

Saul Aryeh Kohn

A DISSERTATION

in

Physics and Astronomy

Presented to the Faculties of the University of Pennsylvania

in Partial Fulfillment of the Requirements for the Degree of Doctor of Philosophy

2018

Supervisor of Dissertation

Graduate Group Chairperson

James E. Aguirre
Associate Professor of Physics and Astronomy

Joshua Klein
Professor of Physics and Astronomy

Dissertation Committee:

Cullen Blake, Assistant Professor of Physics and Astronomy

Adam Lidz, Associate Professor of Physics and Astronomy

Elliot Lipeles, Associate Professor of Physics and Astronomy

Masao Sako, Associate Professor of Physics and Astronomy

to my grandparents

**OUTER SPACE AND FOURIER SPACE:
UNDERSTANDING FOREGROUNDS FOR HI EPOCH OF REIONIZATION
MEASUREMENTS**

© COPYRIGHT

2018

Saul Aryeh Kohn

This work is licensed under the
Creative Commons Attribution
NonCommercial-ShareAlike 3.0
License

To view a copy of this license, visit

<http://creativecommons.org/licenses/by-nc-sa/3.0/>

Acknowledgments

Acknowledgements require a certain mindset to be written well.

ABSTRACT
OUTER SPACE AND FOURIER SPACE:
UNDERSTANDING FOREGROUNDS FOR HI EPOCH OF REIONIZATION
MEASUREMENTS

Saul A. Kohn

James E. Aguirre

Abstracts are written last.

Contents

Title	i
Dedication	ii
Acknowledgments	iv
Abstract	v
List of Tables	x
List of Figures	xi
I Radiation, Interferometry & Polarimetry	1
1 The Epoch of Reionization	3
1.1 This thesis	3
2 Astrophysical Radiation	5
3 Interferometry, Calibration & Polarimetry	6
3.1 The Visibility Equation	6
3.1.1 The Classical Visibility Equation	7
3.1.2 The Measurement Equation	9
3.2 Calibration Techniques	9
3.3 Instrumental Polarization	9
4 Instruments	10
4.1 Instruments used in this work	10
4.1.1 The Donald C. Backer Precision Array for Probing the Epoch of Reionization (PAPER)	11
4.1.1.1 The PAPER signal chain	12
4.1.1.2 PAPER-32	14
4.1.1.3 PAPER-64	15

4.1.1.4	PAPER-128	16
4.1.2	The Hydrogen Epoch of Reionization Array (HERA)	17
4.1.2.1	HERA-19 Commissioning Array	18
4.1.2.2	Future HERA Build-Outs	19
4.2	Other current and future interferometers	20
4.2.1	The Low Frequency Array (LOFAR)	20
4.2.2	The Murchison Widefield Array (MWA)	21
4.2.3	Square Kilometer Array – Low band (SKA-Low)	21
II	Structure in Fourier space	23
5	Peering through the EoR Window	25
6	Data Preparation and Processing	26
6.1	Data Compression	26
6.1.1	Delay–Delay-Rate Filtering	27
6.1.2	Software Implementation	29
6.2	Radio Frequency Interference	31
6.2.1	PAPER-128	31
6.2.1.1	Average Properties	32
6.2.1.2	Individual Properties	39
6.2.1.3	Discussion	42
6.2.2	HERA-19 and PAPER-19	43
6.2.2.1	HERA Hex RFI	44
6.2.2.2	PAPER Hex RFI	46
6.2.2.3	Hex-to-Hex Comparisons	46
6.2.2.4	Comparison PAPER-128 stacked flags	48
6.2.2.5	Discussion	49
6.3	Quality Assurance Metrics	50
6.3.1	Mean Amplitude Flagging	51
6.3.2	Flagging on omnical χ^2	53
7	Polarimetric Calibration	55
7.1	Redundant Calibration	56
7.1.1	Mathematical Overview	58
7.1.1.1	Redundant calibration with OMNICAL	58
7.1.1.2	Including polarization in redundant calibration	59
7.1.1.3	Degeneracies	60
7.1.1.4	Expectations from simulation	61
7.1.2	Data Processing	62
7.1.3	Results	66
7.1.3.1	Calibration	66

7.2	Imaging Calibration	67
8	The Ionosphere	70
8.1	Historical measurements of TEC and RMs	72
8.2	Low frequency observations: discoveries and challenges	74
8.3	Relevance for PAPER and HERA EoR measurements	75
9	A view of the EoR window from the PAPER-32 imaging array	80
9.1	Observations & Reduction	81
9.1.1	Calibration	83
9.1.1.1	Initial calibration	83
9.1.1.2	Absolute Calibration	83
9.1.1.3	Polarimetric factors	85
9.1.2	Creating power spectra	87
9.2	Results	88
9.3	Discussion and Conclusions	91
10	A view of the EoR window from the HERA-19 commissioning array	99
10.1	Polarization Leakage Simulations	100
10.2	Observations & Reduction	100
10.2.1	Calibration	101
10.2.2	Forming power spectra	107
10.3	Results & Discussion	108
10.3.1	General features of the power spectra	109
10.3.2	Day-to-day variability	110
10.3.3	Polarimetric results	111
10.4	Conclusions	113
11	Deep integrations on polarization with PAPER-128	118
11.1	Observations & Reduction	118
11.1.1	Overview of PAPER-128 observations	119
11.1.2	Reduction of Season 1 Epoch 2 data	123
11.2	Results	123
11.3	Discussion & Conclusions	123
III	Expanding the potential of EoR measurements	124
12	Time-Averaged Visibilities	125
13	Higher-order correlation functions between kSZ and 21cm observations	126
14	Deep Learning for 21cm Observations	127

15 Conclusions	128
Appendices	129
A Software	130
A.1 Astronomical Interferometry in Python (<code>aipy</code>)	130
A.2 Astronomy in Python (<code>astropy</code>)	131
A.3 Common Astronomy Software Applications (CASA)	131
A.4 Deep Learning packages	131
A.5 Hierarchical Equal Area isoLatitude Pixelization of the sphere (HEALPix)	131
A.6 <code>pyuvdata</code>	132
A.7 The Scientific Python Ecosystem (<code>scipy</code>)	132
Bibliography	133

List of Tables

6.1	PAPER-128 RFI frequencies and brief characterization for the averaged flags.	34
6.2	RFI as flagged by HERA	45
6.2	RFI as flagged by HERA	46
6.3	RFI as flagged by the PAPER Hex	47
8.1	Ionospheric Layers	71
11.1	PAPER-128 Observing Seasons & Epochs	119

List of Figures

3.1	The geometry of a two-element interferometer, with a plane wave incident from direction \hat{s}	7
4.1	An image of a PAPER dipole element with its groundscreen.	13
4.2	A diagram of the PAPER-128 signal chain. Square-dotted lines indicate RFI-shielding.	14
4.3	The array layouts of the PAPER-32 element deployment in South Africa. .	15
4.4	The array layouts of the PAPER-64 deployment.	16
4.5	The PAPER-128 array layout.	17
4.6	Future vision and physical commissioning HERA layout.	18
4.7	Positions of antennae in the HERA-19 Commissioning Array (including the experimental “subarrays”).	19
4.8	Array sensitivity as a function of k -mode relative to HERA-19	20
6.1	The schema of the database used to organize and implement PAPER data compression.	30
6.2	A waterfall plot of RFI flags averaged over 150 days of PAPER-128 data.	35
6.3	The percentage of time that each frequency was flagged over the season. .	36
6.4	Possible FM radio contamination.	37
6.5	Flights from Cape Town to Johannesburg correspond to RFI in the 120.15 ± 0.35 MHz channels.	38
6.6	The temporal profile of the 5 RFI frequencies with unidentified causes. . .	40
6.7	Waterfalls of RFI flags for nights 2456732, 2456958 and 2457038.	41
6.8	Waterfalls of RFI flags for nights 2456898, 2456924 and 2456965	42
6.9	Frequency vs. percentage flagging for the HERA Hex and PAPER Hex. .	48
6.10	RFI flag waterfalls of frequency vs. South Africa Standard Time for the HERA Hex and PAPER Hex.	49
6.11	An example of the mean amplitude flagging metrics on exercised on HERA-47 data.	52
6.12	Flagging Julian Dates based on the median omnical χ^2 values in the 150–180 MHz band.	54
7.1	Simulations of pseudo-Stokes visibilities,	63
7.2	Arrangement of the PAPER-128 array.	64

7.3	The phase of the complex gain value for the <i>4pol</i> calibration scheme.	67
7.4	Histograms of differences in the ‘x’ gain calibrations per antenna between calibration schemes.	68
7.5	w^2 values, summed across the array, for the <i>x</i> (above) and <i>y</i> (below) feeds in different calibration schemes. From left to right: <i>2pol</i> , <i>4pol</i> and <i>4pol+minV</i> . The color axis is logarithmic, in arbitrary data units. The white gaps are due to RFI flagging.	69
8.1	Distribution of zenithal ionospheric RMs for 3 LSTs in the PAPER-32 observing season	76
8.2	An example of widefield ionospheric RMs calculated by <i>radionoppy</i>	78
8.3	The RM of Cas A as viewed from the LOFAR Core site in the Netherlands on April 11th, 2011, according to <i>ionFR</i> and <i>radionoppy</i>	79
9.1	The PAPER-32, dual-pol antenna imaging configuration and <i>uv</i> distribution.	82
9.2	Snapshot images of Stokes parameters before and absolute calibration. . .	93
9.3	The values of ionospheric RM for different lines of sight a range of LSTs.	94
9.4	The absolute value of delay-transformed visibilities over the bandwidth (146–166 MHz) used to create the power spectra shown in this Chapter. .	95
9.5	Log-scaled 2D power spectra from PAPER-32.	96
9.6	Wedge power spectra from PAPER-32.	97
9.7	Average power in $0.093 < k_{\perp} < 0.098 \text{ hMpc}^{-1}$ as a function of k_{\parallel} for each polarization of PAPER-32 data.	98
10.1	The perimeter of each dish in the HERA-19 array.	101
10.2	Fractional RFI flag occupancy per time and frequency over the eight days of observations.	102
10.3	Multi-frequency synthesis pseudo-Stokes images formed from simulation and data.	104
10.4	Bandpass solutions for the North-South dipole orientation obtained for the functioning antennae in the array on JD 2457548.	105
10.5	The effect of calibration on the phases of visibilities from three redundantly-spaced 14.7 m baselines.	106
10.6	Power spectrum results from the high-band (157–167 MHz).	114
10.7	Power spectrum results from the low-band (120–130 MHz).	115
10.8	High band power as a function of baseline length for the center of the pitchfork ($k_{\parallel} = 0 \text{ h/Mpc}$) and in the EoR window ($k_{\parallel} = 0.2 \text{ h/Mpc}$) for each JD of observation.	116
10.9	Simulated and observed power as a function of k_{\parallel} for the shortest baseline (14.7 m).	117
11.1	An annotated photograph of the PAPER-128 array.	120
11.2	The challenge of Epoch changes.	122

11.3 Good, bad and rotated antennas in Season 1 Epoch 2. 123

Part I

Radiation, Interferometry &

Polarimetry

... His light is like a niche, within which is a lamp, the lamp within glass, the
glass as if it were a white star... Light upon light.

Al Quran, Surah An Nur (24:35)

Chapter 1

The Epoch of Reionization

1.1 This thesis

Everything in this work – algorithmic development, mathematical theory, observations – was carried-out in order to facilitate the detection of the EoR. While these efforts took many forms, they shared that singular motivation of moving the field forward towards a detection of HI at cosmological distances.

This thesis is divided into three parts. Part I is devoted to introducing concepts used throughout this work and building a mathematical formalism around those concepts. Chapter 2 reviews astrophysical mechanisms for producing polarized and unpolarized radiation at low radio frequencies. Chapter 3 builds a formalism around measuring low frequency radio waves with interferometers (and the challenges associated with accurately measuring polarized radiation), and Chapter 4 introduces the instruments used throughout this work.

In Part II I present the bulk of my efforts: building an understanding of the imprint of the polarized sky, and the instrument itself, in the Fourier space used to set limits on the EoR power spectrum. Chapter 5 reviews the current theory and major results of mapping low frequency interferometric measurements into Fourier space. Chapter 6 details several required quality assurance and compression steps that must be taken to clean and interact

with the data. Building from clean data, Chapter 7 presents new algorithms developed to calibrate the measurements. Chapter 8 discusses the impact of Earth’s ionosphere on our measurements. In Chapters 9, 10 and 11 I present successively-deeper integrations on polarized foregrounds in Fourier space.

Part III explores other uses of EoR measurements, beyond detection of the power spectrum. In Chapter 12, I discuss the potential of using long time-averages of interferometric measurements to measure some component of the monopole moment of the sky. In Chapter 13, I present a new formalism for cross-correlating 21 cm emission and CMB anisotropies in Fourier space. Chapter 14 describes my initial investigations into utilizing deep learning techniques for recovering cosmological parameters from simulated EoR measurements. I conclude in Chapter 15.

Chapter 2

Astrophysical Radiation

Chapter 3

Interferometry, Calibration & Polarimetry

In this Chapter I wished to build a formalism around wide-field, polarized interferometric measurements that could be used throughout this work. Many traditional assumptions used in radio interferometry are broken in the case of the wide-field, fully-polarized, drift-scanning measurements native to interferometric EoR observations. In Section 3.1, I derive the equation describing the fundamental observable for an interferometer, called a “visibility”. Section 3.2, I describe calibration techniques relevant to this work and in Section 3.3 I review some of the implications of the previous two sections for polarized measurements.

For a comprehensive review of interferometry from a more traditional perspective, see Thompson et al. (2017).

3.1 The Visibility Equation

A radio interferometer (a term used interchangeably with “interferometric array” for radio observations) is an ensemble of receiving elements, where each element’s measurement is correlated with every other element’s. The simplest case is a two-element interferometer,




Figure 3.1: The geometry of a two-element interferometer, with a plane wave incident from direction \hat{s} .

which we will focus on below. We assume (for now) that the elements are coplanar and identical.

3.1.1 The Classical Visibility Equation

Consider two receiving elements i and j , separated by baseline vector \vec{b} . Suppose a plane wave of wavelength λ is incident upon these elements, with direction of propagation $-\hat{s}$. The geometry of this interferometer is illustrated in Figure 3.1.

We can define the electromagnetic wave to have a frequency dependent phase, such that the electric field measured by element i at time t is

$$E_i = E_0 e^{-2\pi i v t}. \quad (3.1)$$

The time difference between the arrival at i and j is called the “geometrical delay”, τ_g :

$$\tau_g = \frac{\vec{b} \cdot \hat{s}}{c} \quad (3.2)$$

and the electric field measured by element j is

$$E_j = E_0 e^{-2\pi i v(t+\tau_g)} \quad (3.3)$$

An interferometer is an instrument which correlates these electric fields together, integrating their product over some coherent time-scale. This correlation grants:

$$\langle E_i E_j^* \rangle = \lim_{T \rightarrow \infty} \frac{1}{2T} \int_{-T}^T E_i(t) E_j(t) dt = |E_0|^2 e^{-2\pi i v \tau_g} \quad (3.4)$$

where $e^{-2\pi i v \tau_g} = e^{-2\pi i v \vec{b} \cdot \hat{s} / c}$ is known as the “fringe” term, due to its sinusoidal nature. We can generalize this relationship to include more than a single plane wave from direction \hat{s} . Many plane waves, from all directions, can be incident upon the interferometer at a given time and frequency. We can represent the power distribution on the sky as $S(\Omega)$, where $S(\Omega)$. However, no instrument is equally sensitive to radiation from every direction $\hat{s} \in \Omega$. Instead, an instrument has some sensitivity pattern – a *beam pattern* – that tapers the power distribution on the sky into an “observed sky”, $S'(\Omega) = A(\Omega)S(\Omega)$.

These generalizations lead to the classical visibility equation:

$$V_{ij}(v) = \int A(\Omega, v) S(\Omega, v) e^{-2\pi i v \vec{b} \cdot \hat{s} / c} d\Omega \quad (3.5)$$

If we choose to represent the source direction in terms of directional cosines ℓ and m , and represent the baseline vector in units of wavelengths, $\vec{b}/\lambda = (u, v, w)$, we can perform a change of variables in Equation 3.5 to give

$$V_{ij}(u, v) = \int \int A(\ell, m) S(\ell, m) e^{-2\pi i (u\ell + vm + w\sqrt{1-\ell^2-m^2})} d\ell dm \quad (3.6)$$

This relationship is often simplified by assuming only a small area of the sky is under observation – that is, that $A(\ell, m)$ falls-off steeply from zenith – and therefore ℓ^2 and m^2 are small. This grants

$$V_{ij}(u, v) \approx e^{-2\pi i w} \int \int A(\ell, m) S(\ell, m) e^{-2\pi i (u\ell + vm)} d\ell dm \quad (3.7)$$

which plainly casts $V(u, v)$ as the Fourier transform of the observed sky if w is small. Modern low frequency interferometers used in this work greatly violate this approximation, the consequences of which I will discuss in the proceeding sections.

3.1.2 The Measurement Equation

3.2 Calibration Techniques

3.3 Instrumental Polarization

Chapter 4

Instruments

In the following chapters I present data and results from a variety of configurations of two massively redundant low frequency interferometers, PAPER and HERA. In this Chapter I describe these instruments (Section 4.1), along with other current and future low frequency interferometers contributing to EoR science (Section 4.2).

4.1 Instruments used in this work

The vision of Hydrogen Epoch of Reionization Arrays was first laid out in the Backer et al. (2010) White Paper. That work proposed three consecutive efforts, improving upon their predecessors, to construct low frequency interferometers capable of detecting the EoR. While the physical feeds and elements of low frequency interferometers were relatively simple to construct, signal processing, calibration and imaging required new hardware and software to be invented. A research community of observational cosmologists interested in cosmological HI had to be nurtured.

The first of the three stages of Reionization Arrays was a parallel effort. The Precision Array for Probing the Epoch of Reionization (PAPER; Section 4.1.1) and the Murchinson Widefield Array (MWA; Section 4.2.2) investigated separate approaches to¹ antenna

¹Among other things; see Section III B of Backer et al. (2010) for an enumerated list.

design, array layout and calibration techniques, with the objective of setting upper limits on and perhaps detecting the power spectrum of the EoR.

The second stage of the Reionization Arrays brought together the teams from the first stage to design and construct a new interferometer based on the lessons learned from PAPER and the MWA. This new instrument, named *the* Hydrogen Epoch of Reionization Array (HERA; Section 4.1.2) is currently under construction with a build-out schedule that brings new antennas online as they are commissioned. HERA’s objective is not only the detection of the EoR power spectrum, but its characterization at very high signal-to-noise. Attempts at low-fidelity imaging of ionized bubbles will be made.

The nature of the third stage is, at the time of writing, somewhat undetermined and contingent on the next decade of funding for low frequency radio astronomy. In the vision of Backer et al. (2010), its objective will be to image structure evolution throughout the EoR.

4.1.1 The Donald C. Backer Precision Array for Probing the Epoch of Reionization (PAPER)

Much of this thesis presents data from PAPER. PAPER was planned, like HERA, as a staged build-out to larger and larger arrays. In each build-out the correlator was replaced, and (nominally) identical antennae and signal chains were added to the existing array. The first iterations of PAPER consisted of 8 dipole antennae in Green Bank, West Virginia and 4 in Western Australia (Parsons et al., 2010). While the Australian site had a far better observing environment in terms of human-generated radio interference, the site in the USA was easier for the team to design and test on. Green Bank was intended as a test site – brighter foregrounds in the Northern Hemisphere and an unprotected low frequency band inhibited the science goals of the experiment. Nonetheless, the array in Green Bank was build-up to 32 antennae and reconfigured from a more traditional “imaging” configuration to a redundant grid, in order to experiment with redundant calibration and increased sensitivity to discrete Fourier modes (Parsons et al., 2012a; Pober et al.,

2012). Simultaneously, the PAPER-32 array and correlator were constructed in the Karoo Radio Quiet Zone (KRQZ) in South Africa.

4.1.1.1 The PAPER signal chain

The PAPER signal chain changed little throughout the PAPER build-outs in South Africa, and it remained in operation for the HERA-19 Commissioning Array out to HERA-127. HERA elements are actually PAPER feeds, turned upside-down and suspended over a 14 m dish. This heritage was important to understand when interpreting PAPER or HERA data. We briefly describe the PAPER signal chain below. For a more thorough description, refer to Parsons et al. (2010).

Radio waves were incident upon, and induced a voltage in, a dual-polarization PAPER feed. The feed was a sleeved copper dipole protected by a wire-mesh groundscreen. The sleeve broadened the frequency response of the dipole element, and the groundscreen was used to increase sensitivity to emission from zenith (see Figure 4.1 for a photograph).

Electronics next to the dipole element amplified the voltage by a factor of 10^6 , which then propagated down a 150 foot $75\ \Omega$ coaxial cable. All cables were of the same length to minimize the amount of extra calibration required per feed, and were above-ground. These cables ran to 8 “receiverators”; RFI-shielded mini-fridges which contained amplifiers which re-amplified the voltage signals by a factor of 10^4 (to correct for signal loss along the 150 ft coaxial cables) and applied an analog bandpass filter. The filter was designed to have a smooth frequency response which was relatively flat between 120 and 180 MHz (e.g. Moore, 2014).

More 50 foot, $75\ \Omega$ foot coaxial cables ran from these to an RFI-shielded enclosure for further processing. For PAPER and the HERA-19 commissioning array, this enclosure was a specialized shipping container next to the array. For future HERA build-outs, processing will occur in the Karro Array Processing Building (KAPB) and the receiverator architecture will be replaced with underground “nodes” (see DeBoer et al. (2017) for more detail).



Figure 4.1: An image of a PAPER dipole element with its groundscren.




Figure 4.2: A diagram of the PAPER-128 signal chain. Square-dotted lines indicate RFI-shielding.

The filtered analog signal was then digitized with a sampling rate of 100 MHz and passed through an F-engine which Fourier transforms the signal using a 4-tap polyphase filter bank (which allowed for a smooth frequency and Fourier-space response; Price (2016b)). The integration time of each Fourier transform was 10 s. The Fourier transformed signals were distributed over a 10 Gigabit Ethernet switch to the X-engine (Parsons et al., 2008), which cross-multiplied all signals with each other to form visibilities, storing them in MIRIAD files.

A summary of the system described above is shown in Figure 4.2.

4.1.1.2 PAPER-32

The PAPER-32 array in South Africa used a highly redundant configuration in order to take the measurements resulting in, at the time, the strongest upper limits on the EoR




Figure 4.3: The array layouts of the PAPER-32 element deployment in South Africa. *Left:* The redundant grid. Four rows, with each element 30 m from the next in the East-West direction, and closely-packed (~ 4 m) in the North-South direction. Figure taken from Parsons et al. (2014). *Right:* The polarized imaging array. Elements were arranged in a pseudo-random scatter.

power spectrum (Parsons et al., 2014; Jacobs et al., 2015; Moore et al., 2017, sections of Moore et al. 2017 are presented in Chapter 8). The array was in redundant configuration from December 2011 to February 2012. For three nights in September 2011, the 32 elements were reconfigured into an polarized imaging configuration. The results from this deployment were used to make the first 2D power spectra of polarization, presented in Kohn et al. (2016) and in Chapter 9. For images and a brief description, see Figure 4.3.

4.1.1.3 PAPER-64

During the PAPER-32 EoR integration, there were actually 64 antennas present in the Karoo, South Africa. However, the correlator at that time could only process 64 voltage streams – enough for 32 dual-polarization antennae. This is why 64 element single-instrument-polarization imaging results were published prior to any PAPER-32 studies (Jacobs et al., 2013; Stefan et al., 2013). EoR integrations in the 64 element redundant




Figure 4.4: The array layouts of the PAPER-64 deployment. *Left:* The redundant grid, built-out from the PAPER-32 grid. Figure taken from Ali et al. (2015), which highlights the baseline-types used for power spectrum measurements. *Right:* the imaging array, used by Jacobs et al. (2013) to set an absolute flux scale for PAPER experiments. Figure taken from Jacobs et al. (2013).

configuration were only possible after a new correlator was produced in 2012. Results from this array were published in (Ali et al., 2015; Pober et al., 2015, Cheng et al. *in prep.*, Kolopanis et al. *in prep.*). Diagnostic results from the PAPER-64 redundant configuration that informed studies of time-averaged visibilities are presented in Chapter 12. The array layouts are shown in Figure 4.4.

4.1.1.4 PAPER-128

The culmination of the PAPER experiment was the 128 element deployment. There were two observing seasons recorded: November 2013 to March 2014, and July 2014 to January 2015. In this configuration, 112 antennas were laid-out in a redundant grid with 15 m East-West spacings and 4 m North-South spacings. The remaining 16 antennas were arranged in ‘out-rigger’ and ‘in-rigger’ positions to increase *uv*-coverage and enable some




Figure 4.5: The PAPER-128 array layout. 112 antennas were laid-out in a redundant grid with 15 m East-West spacings and 4 m North-South spacings, and the remaining 16 antennas were arranged in ‘out-rigger’ and ‘in-rigger’ positions to increase uv -coverage.

level of imaging. Results from the first observing season of this array are presented in Chapters 6, 7 and 11. The array layout is shown in Figure 4.5.

4.1.2 The Hydrogen Epoch of Reionization Array (HERA)

HERA was ranked as the “top priority in the Radio, Millimeter, and Sub-millimeter category of recommended new facilities for mid-scale funding” by the National Research Council Decadal Survey in astronomy and astrophysics (NRC, 2010). It brought together the largely US-based experts working on PAPER and the MWA to construct a new low-frequency array on the PAPER site in South Africa. The array, comprised of 14 m diameter dishes, was designed to be build-out in stages of close-packed hexagons of increasing size. For a full description of the instrument, refer to DeBoer et al. (2017).

The design of the HERA dish, feed, antenna layout and signal chain are intimately related to lessons learned by the PAPER and MWA EoR teams about the nature of EoR measurements (e.g. Thyagarajan et al., 2015b). One of the major considerations was



Figure 4.6: *Left:* a rendering of the 320 element HERA core. *Right:* the 19 element commissioning array (with the construction team). Leftover PAPER dipoles can also be seen in the background, forming three experimental arrays (described in Section 4.1.2.1). Figure taken from DeBoer et al. (2017)

how the instrument couples to the bright foregrounds, and how to control that coupling – the paradigm of the “wedge” and the “EoR Window”, which are discussed in detail in Chapter 5. Measurements based on prototype feeds were presented in (Ewall-Wice et al., 2016a; Neben et al., 2016, Patra et al. *submitted*).

HERA is a staged experiment, building-out in close-packed hexagons from a 19 element commissioning array, to 37, 127, 240 and finally 350 elements (320 in a dense, fractured core and 30 out-riggers, see Dillon & Parsons 2016 for more detail). Figure 4.6 shows a rendering of the core alongside an image of the HERA-19 commissioning array.

4.1.2.1 HERA-19 Commissioning Array

In October 2015 a HERA commissioning array was completed, connected to the PAPER-128 256-input correlator. This array comprised of four separate components: the first 19 HERA dishes in a close-packed hexagon (HERA-19), 19 PAPER dipoles at the central locations of a hexagon of future HERA dishes (the “PAPER-Hex”), 40 PAPER dipoles in a redundant grid where every-other dipole was rotated 45° (to experiment with different polarization bases; “PAPER Pol”), and the remaining inputs filled with PAPER dipoles arranged in a pseudo-random scatter for imaging (“PAPER-Img”). Studies with this array are presented in Chapter 6 and 10. A diagram of the commissioning setup is shown in




Figure 4.7: Positions of antennae in the HERA-19 Commissioning Array (including the experimental “subarrays”).

Figure 4.7.

4.1.2.2 Future HERA Build-Outs

Each HERA element is much more sensitive than a PAPER element, based purely on collecting area. Figure 4.8 illustrates the varying sensitivities and collecting areas of different arrays and their elements, respectively. PAPER-128 was forecast to make marginal detections of the EoR power spectrum after ~ 1000 hours of integration – HERA-127 should be capable of characterizing the EoR power spectrum at high significance, and the full HERA-350 array will be an extremely powerful survey instrument (e.g. Pober, 2015).




Figure 4.8: Instantaneous array Sensitivities as a function of k -mode (e.g. Chapter 5) relative to HERA-19. Differing element collecting areas are shown on the right. Taken from DeBoer et al. (2017).

We present forecasts for this future instrument in Chapter 13.

4.2 Other current and future interferometers

PAPER and HERA were not the only instruments researching the EoR and how to detect it. Several low-frequency interferometers around the world are contributing to the understanding of the EoR and the difficulties of observing it. Below I briefly describe a few of the leaders of the field – but it is not an exhaustive list.

4.2.1 The Low Frequency Array (LOFAR)

The Low Frequency ARray (LOFAR) is an interferometer made-up of “stations”, the core of which are arranged in a random scatter near Exloo in the Netherlands (van Haarlem

et al., 2013). LOFAR baselines extend across Europe with stations in Ireland, Sweden, Germany, Poland and other locations. These extremely long baselines can provide LOFAR with exquisite imaging capabilities.

LOFAR has detected diffuse, linear polarization structures that may be related to the Milky Way’s magnetic field and interstellar medium (Jelić et al., 2015). Patil et al. (2017) presented deep limits on the EoR power spectrum from one night of LOFAR data (also see Yatawatta et al. (2013)). A series of publications has investigated polarization in Fourier space (Jelić et al., 2014; Asad et al., 2015, 2016, 2017) – results that we build upon and verify throughout this work.

4.2.2 The Murchinson Widefield Array (MWA)

The Murchinson Widefield Array (MWA), based in Murchison Radio-astronomy Observatory in Western Australia, is composed of 256 “tiles” of 16 feeds each. At the time of writing, 72 tiles are arranged in compact hexagonal configurations in an effort to increase EoR sensitivity while retaining imaging capabilities with the rest of the array. The most complete low-frequency catalog of the southern sky, GLEAM, was presented by Hurley-Walker et al. (2017).

The MWA uses separate analysis pipelines for its EoR measurements (Sullivan et al., 2012a; Jacobs et al., 2016; Trott et al., 2016), a concept that will be implemented on future HERA measurements. The MWA has detected polarized signal from diffuse Galactic emission (Lenc et al., 2016, 2017) and set deep limits on the EoR power spectrum (e.g. Dillon et al., 2014, 2015). Pioneering work on the ionosphere was accomplished using MWA data (e.g. Loi et al., 2015), which we discuss in Chapter 8.

4.2.3 Square Kilometer Array – Low band (SKA-Low)

The Square Kilometer Array (SKA) will be built on two sites: the Karoo Radio Quiet Zone, currently occupied by HERA, will be the central location of the “high–mid band”

(350 MHz – 14 GHz; “SKA-Mid”) observatory and the Murchison Radio-astronomy Observatory, currently occupied by the MWA, will be the central location of the “low band” observatory (50 – 350 MHz; “SKA-Low”).

SKA-Low will consist of over 100,000 receiving elements in an imaging configuration. It will be the most powerful low-frequency radio telescope ever created, and be used not only for EoR science but a host of low-frequency science objectives (e.g. Carilli & Rawlings, 2004; Schilizzi et al., 2007; Dewdney et al., 2009). Learning from the work of EoR teams working on PAPER, HERA, MWA, LOFAR and other instruments will be crucial to the success of SKA-Low’s EoR science goals.

Part II

Structure in Fourier space

The Gemara asks: Do we not see that comets pass Orion? The Gemara rejects this: The aura of the comet passes Orion and it appears as though the comet itself passes. Rav Huna, son of Rav Yehoshua, said a different answer: It is merely that veil, one of the firmaments, rips and furls and the light of the next firmament is seen, and this appears like a comet. Rav Ashi said another explanation: It is not a comet that passes Orion, but a star that is uprooted from one side of Orion, and another star, from the other side of Orion, sees it and gets startled and shudders, and appears as if it is passing.

Talmud Berakhot 58b:11

Chapter 5

Peering through the EoR Window

Chapter 6

Data Preparation and Processing

The data volume of interferometric measurements inherently scale as the square of the number of antennas in the array (N_{ant}). Not only does the sheer volume of data from large- N_{ant} arrays pose a problem for data storage, but also it requires precise and efficient efforts to quality assure (QA) the data.

In this chapter, I will outline some of the efforts involved in data preparation, pre-processing and QA that are required for an EoR power spectrum estimate.

6.1 Data Compression

The PAPER-128 correlator produced 288 MIRIAD files per night. Each of these contained 8126 baselines, and each baseline contained visibilities over 1024 98 kHz frequency channels and 56 10 s time integrations. The four instrumental polarizations were in separate files. In sum, each file was 4.2 GB which meant that each night 1.2 TB of data were recorded.

In order to efficiently transport the data over Gigabit Ethernet from the Karoo Radio Quiet Zone (KRQZ) to Cape Town, and from Cape Town under transatlantic cables to Philadelphia, some compression was required. It was also required that such a compression, while lossy, did not effect the targeted cosmological signal.

6.1.1 Delay–Delay-Rate Filtering

The compression algorithm implemented for PAPER observations, Delay–Delay-Rate (DDR) filtering, was introduced in Parsons & Backer (2009) described in Parsons et al. (2014), and we briefly review it below.

The geometric delay of a celestial signal, originating from direction \hat{s} , incident on an interferometric baseline described by vector \vec{b} , is

$$\tau_g = |\vec{b} \cdot \hat{s}|/c \quad (6.1)$$

where c is the speed of light. This relationship implies that τ_g is bounded for a given baseline

$$-|\vec{b}|/c \leq \tau_g \leq |\vec{b}|/c \quad (6.2)$$

Equation 6.2 therefore gives the maximum value of $|\tau_g|$ physically meaningful for a given array – the maximum baseline length in that array, divided by c . For PAPER, the maximum baseline length is 300 m, corresponding to $\max(|\tau_g|) = 1\mu\text{s}$. As reviewed in Chapter 5, the delay axis may be accessed by Fourier transforming a visibility along the frequency axis. Once in delay space, power at delays larger in magnitude than $1\mu\text{s}$ could be removed. With a sufficiently large frequency bandwidth, this would not produce aliased signal, according to the critical Nyquist rate. By using the $1\mu\text{s}$ as a delay bound for all visibilities, the frequency axes of all compressed visibilities remained the same (reduced in number from 1024 to 203), which while sub-optimal from a compression point of view, allowed for ease of programming at later stages.

A similar geometric bound can be obtained by Fourier transforming the time axis of visibilities, provided that they were obtained in drift-scan mode (see Chapter 3). Parsons & Backer (2009) showed that the rate at which the geometric delay on an interferometric baseline changes is governed only by the position of the array on Earth, and the Earth’s rotation:

$$\dot{\tau}_g = -\frac{\omega_{\oplus} \cos \delta}{c} (b_x \sin \alpha + b_y \cos \alpha) \quad (6.3)$$

where ω_{\oplus} is the angular frequency of the Earth's rotation, α and δ are the hour-angle and declination of a point on the celestial sphere, respectively, and $\vec{b} = (b_x, b_y, b_z)$ is the baseline vector expressed in equatorial coordinates.

For arrays not close to the geographic poles, $|b_y| \gg |b_x|$, there is a maximum rate of change (corresponding to $(\alpha, \delta) = (0, 0)$), producing a bound on $\dot{\tau}_g$:

$$-\omega_{\oplus} |b_y|/c \leq \dot{\tau}_g \leq \omega_{\oplus} |b_y|/c \quad (6.4)$$

for a 300 m East-West baseline, the maximum delay-rate is approximately $\max(|\dot{\tau}_g|) = 0.07 \text{ ns s}^{-1}$. This delay-rate was not Nyquist sampled by a single PAPER file: requiring the previous and next files generated for that polarization to be appended on either side of each visibility's time axis to prevent aliasing from the decimation. For the large scale processing of months of data, this required a software pipeline described in Section 6.1.2.

There are also other issues with DDR compression, largely associated with instrument systematics. Delay transforms rely on the fact that the bright foregrounds that dominate the measured signal are spectrally smooth, and that the frequency response of the instrument is also spectrally smooth: this of course is the basis for the EoR window paradigm reviewed in Chapter 5. Likewise, delay-rate filtering assumes temporal smoothness. Radio Frequency Interference (RFI) signals created by human communications violate both models of smoothness, since they are typically confined to narrow bandwidths (creating sharp spikes along the frequency axis) and may be transient (creating sharp spikes along the time axis). This requires steadfast identification and flagging algorithms for RFI (see Section 6.2), and some variety of interpolation, fitting, or CLEANing across the flagged regions prior to compression.

By DDR filtering of PAPER-128 data using a 300 m baseline to set the width of the filters we were able to reduce the volume of the data by an approximate factor of 70.

6.1.2 Software Implementation

The first season of PAPER-128 data, due to a variety of circumstances, required compression on the computing cluster at the University of Pennsylvania. The raw data were stored on a high-volume drive that was able to connect with the cluster via a low-speed switch. The hardware capable of performing any sort of high-performance processing (i.e. holding the data in RAM) were ten “compute nodes” connected to the cluster via a high-speed switch, and mounted in an NFS architecture. The compute nodes could only hold ~ 10 PAPER-128 files in storage.

The processing stages for compression of a night of PAPER data, described below, required knowledge of the location and compression state of not only individual files, but also the neighbors-in-time of the file in question, in order to implement the DDR filter described above. To supervise the compression we created a MySQL database, which we interacted with via Shell and Python scripts. The database contained a table for the data files under processing and their compression state, a table of neighbor-relations, a table of file details, and a table of the processing nodes available. The schema of this database is shown in Figure 6.1.

To implement the compression, per file, the following steps were required:

1. Copying the file from the storage volume to the cluster. For a single night of data, this required roughly 8 hours.
2. Copying the file from the cluster to the compute node. This required roughly 5 minutes.
3. Generate copy of the file, with metadata corrections. This required roughly 1 minute.
4. Delete the raw file.
5. RFI-flag the high frequency-resolution data. This required roughly 2 minutes.
6. Delete the metadata-corrected file.




Figure 6.1: The schema of the database used to organize and implement PAPER data compression.

7. Acquire time-neighbors to the file in question, and bring them to the RFI-flagged stage. The time required for this stage varied with cluster activity, but usually required roughly 20 minutes.
8. DDR filter the RFI-flagged data, using an high-tolerance iterative CLEAN. This required roughly 20 minutes.
9. RFI flag the compressed data (coarse flagging), saving the flags to a separate file. This required roughly 1 minute.
10. Apply the coarse RFI flags to the *uncompressed*, RFI-flagged data. This required roughly one minute.
11. DDR filter the now twice-RFI-flagged data, using a low-tolerance iterative CLEAN. This required roughly 120 minutes.
12. Copy compressed data to the cluster.

13. Delete the twice-RFI-flagged data.
14. If the once-RFI-flagged data are not required as neighbors, delete them.
15. Delete the compressed data from the compute node.
16. If neighbors have already been compressed, delete them, otherwise begin their compression.
17. After all files are compressed, delete the uncompressed files from the cluster.

In total, this meant that across ten compute nodes, and efficient use of the fact that the neighbors could progress through the processing stages while the central file was being compressed, meant that it took roughly 20 to 24 hours to compress a night of observations.

6.2 Radio Frequency Interference

As noted above, RFI was able to introduce spectral and temporal structure that would cause ringing in the data during compression if it was not flagged. This meant that both identification and characterization of RFI was crucial to the scientific goals of the PAPER and HERA experiments. In Section 6.2.1, I present characterization of RFI in the second season of PAPER-128 data. By averaging flags in local time I was able to investigate “repeat offender” frequency bands and identify outlying “quiet” and “loud” days. In Section 6.2.2 I analyze RFI flags from the first Internal Data Release (IDR1) of HERA commissioning data, which contained 19 HERA feeds suspended above 5 m dishes in a close-packed hexagon, and 19 PAPER feeds in the same positions as the central dishes, allowing us to investigate the difference in flagging between feeds at different altitudes.

6.2.1 PAPER-128

The PAPER-128 2014 observation season ran from 18th June 2014 through the 30th April 2015. During this run, some 150 nights of data were recorded. A “night”, which I will

refer to using the JD at the start of observations, consists of twelve hours of observation from 6pm to 6am South African Standard Time (SAST). Observations as processed by the PAPER correlator are recorded in MIRIAD uv files. These files contain visibilities for each antenna pair in the array. Each integration is 20 seconds long over 1024 frequency bins from 100 to 200 MHz. Each uv file contains 56 integrations per antenna pair, and 72 uv files are recorded per linear polarization (xx, xy, yx, yy) per night.

Early in the PAPER data compression process, visibilities are flagged for RFI. This is accomplished by the aipy script *xrfi_simple.py*, which takes the derivative of the frequency axis of all baselines associated with a single antenna, and flags any frequencies with a derivative $\geq 6\sigma$ above the mean. We always flag the band-edges (~ 7 MHz on each side), since these frequencies are not useful to us, and always flag the 137 ± 0.6 MHz band associated with ORBCOMM satellite network transmissions. This process is repeated per integration within each uv file and stored in a Python numpy zip (npz) file. This means that any baseline associated with antenna 1 can contribute a flag to the resultant npz file, which in turn is applied to the data.

The result is 280 files of high-time and -frequency resolution files per night per linear polarization containing information about the RFI environment of the HERA site. I report on the properties of these flags in time- and frequency-space over the 2014 observation season. This section is organized it as follows: in Section 6.2.1.1, I analyse the average properties of RFI over the season by stacking flags in local time and normalizing appropriately. In Section 6.2.1.2 I address nights with particularly strange RFI properties. I discuss the implications of my findings in Section 6.2.1.3.

6.2.1.1 Average Properties

In order to assess the average properties of the RFI environment, I calculated a weighted average of flags over the season. Over 150 nights, one-time occurrences are washed-out beneath the 1% level, allowing me to assess persistent issues.

Nominally, each night should grant 3920 integrations-worth of flags over 1024 fre-

quency bins, per linear polarization. In reality, most of the time this holds true, but occasionally not all files are compressible (hence failing to generate flags) or observations fail to start at the correct time (so there are no data to flag). Also, in the event of an X-engine failure within the correlator, contiguous chunks of the band (in eighths, i.e. 25 MHz across) are flagged-out, usually for the rest of the night.

For this reason, I calculated a weighted average of the flags across the season, but neglected nights with correlator failures or late starts. Weights were simply the number of nights that contained that integration-bin in SAST. The resultant “flag density waterfall” is shown in Figure 6.2. The color scale is indicative of flagging frequency across the season, and line plots above and to the right of the the waterfall showing the percentage of times and frequencies that were flagged, respectively.

A summary of the persistent (flagged $\geq 1\%$ of the time per channel) RFI frequencies can be found in Table 6.2.1.1. I have investigated each frequency and tried to find the most likely source for each. In most cases, this required looking at the properties in time as well as frequency. Others were more obvious from frequency alone, e.g. the 149.8 MHz transmission frequency from the International Space Station (ISS). Still others I could not track down a convincing explanation for, and these are listed with a ‘?’ . A ‘?’ next to a possible cause indicates that the listed cause is the most prevalent at that frequency, but that the temporal properties of that cause do not necessarily make sense. Many of the characterizations arise from the South African Table of Frequency Allocations (SATFA; Staatskoerant (2008)).

Figure 6.3 shows the detail of the top panel of Figure 6.2. This figure highlights the broad swath of the band from roughly 150 to 180 MHz that was, on average, clear of RFI. This roughly corresponds to 21 cm redshifts $z = 6.9$ to 8.5 . This is one of the reasons that the Parsons et al. (2014) and Ali et al. (2015) limits on the 21 cm power spectrum concentrated on this redshift range – there were simply more unflagged data to average-down with. Furlanetto et al. (2006) show that the $z \sim 8$ universe can be considered roughly coeval over an ~ 8 MHz bandwidth. As such, the 30 MHz chunk could be used to

Table 6.1. PAPER-128 RFI frequencies and brief characterization for the averaged flags.

ν MHz	Flagged %	Cause (Possible)	Notes or Time (SAST) Characterization
103 ± 3	100	BAND EDGE	Built-in to flagger.
107.25 ± 0.25	2.6	FM radio	Constant background at 2% level
107.55 ± 0.05	1.9	FM radio	Constant background at 2% level
108.1 ± 0.4	9	FM radio?	Rises with time, peaking at midnight and 4am
109 ± 0.4	11.5	FM radio?	Rises with time, peaking around 4am
112.8 ± 0.1	1.4	Aircraft?	Constant background at 1% level
114.05 ± 0.85	3.7	?1	Decreases till midnight; peak at 4am
116.55 ± 0.35	2.2	?2	Peak at midnight
120.15 ± 0.35	3.2	Aircraft	Roughly follows CPT↔JNB flight times
124.95 ± 0.35	5.5	Aircraft	Roughly follows CPT↔JNB flight times
130.25 ± 0.55	4.3	?3	Falls (7pm) and rises (3am) steeply
131.75 ± 0.35	10.3	Aircraft?	Peaks at 6:30, 7:30, 8:30, 9:30, 10 and then a steep falloff
136.05 ± 0.45	33.1	Radar?	Decreases over night
137.35 ± 0.85	100	ORBCOMM	
141.45 ± 0.35	2.1	Mobile phones?	High until 9pm, then at background 1% level
145.85 ± 0.45	10.7	Amateur radio	Strong 9pm-1am – this is the official downlink for ISS-HAM
149.75 ± 0.55	90.7	ISS	“Beeps”, but in stacked data peaks 2am
175.15 ± 0.35	20.5	VHF TV (video)	Channel 4. Peaks at 8:30pm, then falls to background 7%
181.15 ± 0.15	1.6	VHF TV (audio)	Channel 4. 2% level turns-off at 10pm
182.15 ± 0.35	75	?4	Decreases until 10pm (to 15%), when it begins a slow rise again
183.2 ± 0.5	89.7	VHF TV (video)	Channel 5. Rises throughout night.
186.25 ± 0.35	4.6	?5	Extreme turn-off at 9:45
189.15 ± 0.35	41.4	VHF TV (audio)	Channel 5. Rises throughout night.
189.9 ± 0.4	100	VHF TV	Channel 6. Built-in to flagger.
191.1 ± 0.3	100	VHF TV	Channel 7. Built-in to flagger.
196 ± 4	100	BAND EDGE	




Figure 6.2: A waterfall plot of RFI flags averaged over 150 days of data. The gridding process is described in the text. Above the waterfall I show the percentage of the season each frequency is flagged, and to the right I show the percentage of frequencies that are flagged per integration.

create ~ 3 power spectra, as demonstrated in Jacobs et al. (2015). As we show below, the deactivation of VHF TV broadcasts could enable measurements up to the band edge.




Figure 6.3: The percentage of time that each frequency was flagged over the season.

FM Radio

SATFA lists the frequency band 87.5–108 MHz as available for FM radio broadcasts, leading me to postulate that the low-level RFI we observed in the 107.25 ± 0.25 and 107.55 ± 0.05 MHz bands had FM radio as the leading cause. The 108.1 ± 0.4 and 109 ± 0.4 MHz bands were outside of the official range, and exhibit odd temporal properties for human activity – two peaks at midnight and 4am – with a increasing number of flags throughout the average night (see Figure 6.4).

Aircraft communications

It was difficult to argue that the 112.1 ± 0.1 MHz signal is caused by aircraft communications since it maintained a constant background level. However, SATFA listed this frequency as reserved for aircraft communications and it has been used in the past as a calibration frequency for aircraft instruments South African Civil Aviation Authority (2008).

The other aircraft frequencies were obvious, because they closely traced the 2-hour




Figure 6.4: Possible FM radio contamination in the *Top, left to right*: 107.25 ± 0.25 and 107.55 ± 0.05 MHz bands, and *Bottom, left to right*: 108.1 ± 0.4 and 109 ± 0.4 MHz bands.

flight from Cape Town to Johannesburg¹. An example (120.15 ± 0.35 MHz) is shown in Figure 6.5. SATFA reserved frequencies 108–117.975 MHz for aeronautical radionavigation and 117.975–137 MHz for aeronautical mobile. In Table 6.2.1.1 I listed 131.75 ± 0.35 MHz as caused by aircraft since it falls in the aeronautical mobile band, but it does not follow the flight patterns as closely as the other bands.

Orbital communications

ORBCOMM Inc.’s constellation of 29 LEO communication satellites is a well-known contaminant of the low-frequency sky, dominating over any astronomical signal at 137–138 MHz (although each satellite emits within a 20 kHz band). For this reason there was

¹Credit to Danny Jacobs for first spotting this and noting it in an internal PAPER circular in December 2009.




Figure 6.5: Flights from Cape Town to Johannesburg correspond to RFI in the 120.15 ± 0.35 MHz channels. Vertical dashed lines indicate a flight leaving Cape Town (flights from Johannesburg are roughly concurrent) and the flight code is listed. The transparency of a line is inversely proportional to how many days a week that flight is scheduled for. The flight is 2 to 2.5 hours long – and about 2 hours after the last flight of the day, the flags fall to background level (but notably, not always to zero).

built-in flagging at 137.35 ± 0.85 MHz within the compression pipeline.

The largest contaminant without built-in flags in the pipeline were communications from the ISS. The 149.75 ± 0.55 MHz transmissions were semi-regular in time; they ‘beep’.

Onboard the ISS are HAM radio devices. Some countries have also launched satellites with these onboard, one of the purposes of which is to provide HAM radio operators something in space to communicate with. These devices are licensed to operate at 145.2 and 145.8 MHz, and SATFA listed the 144-146 MHz band as reserved for ‘Amateur–Satellite’ communications. We detected RFI at 145.85 ± 0.45 MHz, although strong signal across $\sim 10\%$ of the season that occurs 9pm-1am argues against human operation.

Mobile phones and VHF TV

A weak RFI signal at 141.45 ± 0.35 MHz was within the ‘mobile 1 BTX’ and aeronautical mobile band in SATFA, but other than this single listing I did not build a strong case for

the signal's cause.

VHF TV is broadcast over specifically-spaced video and audio frequencies. The strong signals at 183.2 ± 0.5 MHz and 189.15 ± 0.35 MHz had almost identical gradients for the percentage of flagging as a function of time of night. These frequencies corresponded exactly to Channel 5 of South African System I 625-line VHF TV signals for video and audio transmission, respectively. Similarly, the weaker signals at 175.15 ± 0.35 and 181.15 ± 0.15 MHz corresponded to Channel 4's video and audio transmission, respectively, but they did not share the same temporal properties.

Unidentified sources

There were 5 RFI frequencies in the averaged data that I could not identify the sources of: weak emissions (flagged < 5% of the season) at 114.05 ± 0.85 , 116.55 ± 0.35 , 130.25 ± 0.55 and 186.25 ± 0.35 MHz, and one strong emission at 182.15 ± 0.35 MHz. The variation of each source with time is shown in Figure 6.6. The 186.25 ± 0.35 MHz had a sharp turn-off around 9:45pm each night, suggesting that it originated from some kind of automated device.

6.2.1.2 Individual Properties

Using the flags per night, I was able to assess the total number of flags as a percentage of the waterfall (i.e. $N_{\text{flags}} / (3920 \times 1024)$). The average flagging per night was $19.2 \pm 0.5\%$, which was dominated by the permanent flagging of ORBCOMM and band edges. Four nights deviated from the average by a $\geq 2\sigma$ excess: JDs 2456965, 2456732, 2456958 and 2457038. Their flag waterfalls are shown in Figure 6.7 (2456732, 2456958 and 2457038) and Figure 6.8 (2456965). While the strange nature of night 2456965 is discussed below, the three others followed the pattern of having strong contamination from FM and aircraft communication bands, but also had broadband 'pulses' up to about 20 minutes in length. The source of these broadband pulses is not well understood, although it was clear that




Figure 6.6: The temporal profile of the 5 RFI frequencies with unidentified causes. *Top, left to right:* 114.05 ± 0.85 and 116.55 ± 0.35 MHz. *Middle, left to right:* 130.25 ± 0.55 and 182.15 ± 0.35 MHz. *Bottom:* 186.25 ± 0.35 MHz. The 182.15 ± 0.35 MHz frequency is flagged a large amount of the time, making it our most-offending unidentified source.

ORBCOMM tends to spill outside of its allocated band on occasion.

JD 2456965 was easily the worst offender, and it exhibited a strange signal that wanders in frequency and time close to the ISS band. An event of note on this date (23rd November 2014) was a Soyuz FG launch that docked with the ISS – this may have been




Figure 6.7: *Left to right:* waterfalls of flags for nights 2456732, 2456958 and 2457038. These three nights, along with 2456965, are $>20.2\%$ flagged; $>2\sigma$ above the average flagging amount per night.

a signature of their transmissions². Similar signals were seen on 2456898 (28th August 2014; although only at the beginning of the night) and 2456924 (23rd September 2014). There was no listed orbital or suborbital activity for 2456898. There was an US ICBM test off of the coast of Virginia on 2456924, but this was probably not the cause of the RFI. The flag waterfalls for these nights are shown in Figure 6.8.

Another property that the flag waterfalls in Figures 6.7 and 6.8 highlight is the presence of broadband RFI signals, typically present at frequencies lower than the ORB-COMM band. However, while we flagged at the low-end of the band (which had higher noise levels to begin with), it is likely that such broadband pulses dominated the band at those times, and that we failed to flag all of the integrations. Our flagging routine `xrfi_simple.py` does contain a thresholding option for flagging the entire integration given some arbitrary number of frequencies flagged during that integration: some experimentation will be required to decide if that threshold should change.

²The internet also suggests... less plausible explanations: <https://www.youtube.com/watch?t=11&v=VtZx8iP04zs>.




Figure 6.8: *Left to right:* waterfalls of flags for nights 2456898, 2456924 and 2456965. These three nights exhibit a strange behaviour of RFI that changes in frequency and time. JD 2456965 is by far the worst, and during this night as well as 2456898, we see a broadband ‘comb’ of flagged frequencies near the band edges

6.2.1.3 Discussion

Based on my findings, I was able to recommend some actions that could be taken in the KRQZ to enable better measurements:

- Steps to reduce and ideally eliminate the VHF TV transmissions in the area would be very helpful, since these were clearly interfering with our measurements in the high-end of the band.
- The ISS 149.75 ± 0.55 MHz band should be permanently flagged within the compression pipeline.
- Pursuing re-routing of flight paths will not do much to help: we see aircraft signals for the duration of their flight, not just when they’re over the Karoo.
- A lower threshold for identifying broadband RFI should be investigated.

A new, lower-frequency feed is currently under development by the HERA analog group. This would nominally allow measurements to be taken in the range 50–250 MHz, allowing science observations of the Dark Ages and the post-reionization Universe. It should be noted that at the lowest frequencies FM radio will be a constant harassment to

these measurements. At the higher frequencies, VHF TV will be the primary contaminant, but should be much easier to remove as it is both narrow-band and within the KRQZ’s power to shut off.

6.2.2 HERA-19 and PAPER-19

The HERA-19 IDR1 consisted of four subarrays: the HERA-19 hexagon of dishes (the ‘HERA Hex’), a hexagon of 19 PAPER dipoles in the exact positions of not-yet-constructed HERA dishes (the ‘PAPER Hex’), an imaging array and an experimental array for polarization measurements. I will concentrate on the two Hexes in this Section. I analyzed RFI as flagged in the linear xx -polarization only. Asymmetric beams can in principle receive different RFI events for different linear polarizations, but analysis of that was outside the scope of this diagnostic study.

IDR1 consisted of one ‘golden day’, JD 2457458. This ran from 6pm on March 10th 2016 to 6am the following day. This gave, per baseline, roughly 4000 integrations of 10 seconds each over 1024, 100 kHz frequency channels from 100 to 200 MHz.

In order to flag RFI I used the aipy script *xrfi_simple.py*. I took the union of all baseline flags as data to analyze. Unlike in Section 6.2.1, these data did not have a-priori flagging of band edges, which allowed me to make a more complete study of RFI in the HERA band. I did have to implement custom flags in order to get more than a zeroth-order view of the RFI (since these would dominate the flagging routine unless they are flagged already), but my results from Section 6.2.1 gave a better idea of what was flagged to get there.

Below I present measurements of high-power, mostly narrow-band RFI channels as flagged in HERA Hex data and PAPER Hex data separately. In both cases, I was able to list any channels that are flagged for $\geq 1\%$ of the night. I could then compare the flagging Hex-to-Hex, and to PAPER-128.

6.2.2.1 HERA Hex RFI

Table 6.2.2.1 shows all narrowband frequency ranges flagged in HERA-19 visibilities, with columns of the frequency range in MHz, % flagging over time, plausible identification, whether or not it was identified in PAPER-128 data, and other notes (often details of the possible identification). Frequencies with 100% flagging indicate manual flags required for `xrfi_simple` to work on the rest of the channels.

Clearly, the low-end of the band was swamped by FM radio broadcasts. One notable frequency was the 109.2 ± 0.3 MHz band, which was heavily flagged in HERA visibilities, but was only flagged a few percent in PAPER-128 data.

As seen before, ORBCOMM satellite emissions spilled out of their allocated 137–138 MHz band down to 136.3 MHz.

There were many narrowband RFI channels, across the band, that PAPER-128 did not pick-up. Most of these were flagged only at low levels, with two exceptions: 111.3 ± 0.2 MHz and 113.5 ± 0.1 MHz. Both of these were in the aircraft navigation band. There is some evidence (Civil Aviation Authority, 2016) that 111.3 MHz band is for air force communications. The 113.5 MHz band is a known band for radionavigation beacons ('VOR navaids' World Aero Data, 2016).

A particularly annoying ‘new’ emitter was in the 153.8 ± 0.2 MHz region, which is close to the center of our nominal EoR band. It could correspond to mobile phones being used close to site.

Table 6.2. RFI as flagged by HERA

ν MHz	Flagged %	Cause (Possible)	Seen by PAPER-128	Notes
100.7 ± 0.2	50	FM Radio	n/a	RSG "Dis Die Een" Prieska
101.5 ± 0.3	36	FM Radio	n/a	RSG "Dis Die Een" Calvinia
102.4 ± 0.1	100	FM Radio	n/a	RSG "Dis Die Een" Carnarvon
102.8 ± 0.3	57	FM Radio	n/a	RSG "Dis Die Een" Pofadder
104.2 ± 0.1	100	FM Radio	n/a	SAfm Prieska
105.1 ± 0.2	100	FM Radio	n/a	SAfm Calvinia
106.2 ± 0.3	100	FM Radio	n/a	SAfm Carnarvon
106.9 ± 0.1	15	FM Radio	n/a	Sentech
107.2 ± 0.1	18	FM Radio	Yes	
107.8 ± 0.2	15	FM Radio	Yes	
108.3 ± 0.1	31	FM Radio?	Yes	
109.2 ± 0.3	93	FM Radio?	Yes...	...but not to this degree
111.3 ± 0.2	25	Air force?	No	
112.5 ± 0.1	5	Aircraft?	No	
113.5 ± 0.1	21	Aircraft	No	VOR navaids
115.5 ± 0.1	3	Navaids?	No	
115.9 ± 0.1	3	Navaids?	No	
116.6 ± 0.2	9	Aircraft?	Yes	VOR-DME navaids
120.1 ± 0.2	5	Aircraft	Yes	CPT< – >JNB
125.0 ± 0.2	6	Aircraft	Yes	CPT< – >JNB
130.0 ± 0.2	4	Aircraft	No	Communication
131.6 ± 0.2	15	Aircraft	Yes	KLM OPS
136.4 ± 0.1	9	ORBCOMM	Yes	
136.7 ± 0.1	10	ORBCOMM	Yes	
137.4 ± 0.4	100	ORBCOMM	Yes	
145.7 ± 0.4	18	ISS/Amateur Radio band	Yes	
149.9 ± 0.1	100	ISS	Yes	
153.8 ± 0.2	7	Mobile phones?	No	
175.0 ± 0.1	100	VHF TV	Yes	Channel 4 Video
178.3 ± 0.2	8	VHF TV	No	Channel 7?
181.2 ± 0.1	100	VHF TV	Yes	Channel 4 Audio
182.2 ± 0.2	9		Yes	
183.5 ± 0.6	100	VHF TV	Yes	Channel 5 Video
184.1 ± 0.1	2	VHF TV?	Yes	Channel 5?
184.7 ± 0.1	6	Broadcasting	No	
187.8 ± 0.1	4		No	
189.1 ± 0.1	52	VHF TV	Yes	Channel 5 Audio
190.1 ± 0.3	13		n/a	

Table 6.2—Continued

ν MHz	Flagged %	Cause (Possible)	Seen by PAPER-128	Notes
191.1 ± 0.1	100	VHF TV	n/a	Channel 7
197.2 ± 0.2	18		n/a	
199.4 ± 0.5	100	BAND EDGE	n/a	

6.2.2.2 PAPER Hex RFI

Table 6.2.2.2 has the same description as Table 6.2.2.1, but for the PAPER Hex. There were far fewer RFI frequencies flagged in PAPER visibilities, almost all of which were seen by HERA. The only RFI seen by the PAPER Hex and not the HERA Hex was the 123.5 ± 0.1 MHz emission, which I could find a plausible identification for.

6.2.2.3 Hex-to-Hex Comparisons

As mentioned above, the PAPER Hex saw far fewer narrowband RFI channels than HERA does. This highlighted an interesting trade-off between dipoles and dishes: at first glance, one might have expected PAPER dipoles to be more susceptible to RFI given their broader effective beams. However, HERA dipoles are lifted several meters above the ground, and this change in height may have been the source of the greater susceptibility to RFI. RFI comes from the horizon, which would be more easily received in the far sidelobes of the beam.

Even for the RFI channels they did share, HERA flagged them more often. Taking the difference in percentage-flagging for the common RFI channels (think of the left panel subtracted from the right panel for common channels in Figure 6.9), those channels had an average of 8% more flagging in HERA visibilities. The difference was particularly high in the aeronautical radionavigation bands, where HERA had on average 38% more flagging than the PAPER Hex.

Figure 6.10 shows the flags on a per-sample basis (these were averaged over time to

Table 6.3. RFI as flagged by the PAPER Hex

ν MHz	Flagged %	Cause (Possible)	Seen by PAPER-128	Notes
100.0 ± 0.1	100	BAND EDGE	n/a	
100.7 ± 0.1	11	FM Radio	n/a	RSG "Dis Die Een" Calvinia
101.6 ± 0.2	6	FM Radio	n/a	RSG "Dis Die Een" Calvinia
102.4 ± 0.1	100	FM Radio	n/a	RSG "Dis Die Een" Carnarvon
102.7 ± 0.1	100	FM Radio	n/a	RSG "Dis Die Een" Pofadder
104.2 ± 0.2	100	FM Radio	n/a	SAfm Prieska
105.1 ± 0.2	100	FM Radio	n/a	SAfm Calvinia
106.2 ± 0.3	100	FM Radio	n/a	SAfm Carnarvon
108.2 ± 0.1	3	FM Radio?	Yes	
109.1 ± 0.1	26	FM Radio?	Yes	
113.6 ± 0.1	2	Airplane Communications	No	VOR navaid
120.2 ± 0.3	3	Aircraft	Yes	CPT< - >JNB
123.5 ± 0.1	1		No	Not seen by HERA
125.0 ± 0.2	6	Aircraft	Yes	CPT< - >JNB
130.0 ± 0.3	3		No	
131.7 ± 0.2	14	Aircraft	Yes	
136.4 ± 0.2	6	ORBCOMM	Yes	
136.7 ± 0.2	6	ORBCOMM	Yes	
137.4 ± 0.4	100	ORBCOMM	Yes	
145.8 ± 0.3	14	ISS/Amateur Radio band	Yes	
149.9 ± 0.1	100	ISS	Yes	
153.8 ± 0.2	3	Single frequency mobile phones?	No	
175.1 ± 0.2	100	VHF TV	Yes	Channel 4 Video
178.3 ± 0.2	100	VHF TV	No	Channel 7?
181.2 ± 0.1	100	VHF TV	Yes	Channel 4 Audio
183.2 ± 0.2	100	VHF TV	Yes	Channel 5 Video
189.2 ± 0.1	100	VHF TV	Yes	Channel 5 Audio
191.2 ± 0.1	100	VHF TV	n/a	Channel 7
199.8 ± 0.2	100	BAND EDGE	n/a	




Figure 6.9: Frequency vs. percentage flagging for the HERA Hex (*left*) and PAPER Hex (*right*). Any band with greater than 1% flagging is reported in Tables 6.2.2.1 and 6.2.2.2.

create Figure 6.9). Most apparent was the occupancy of the HERA plot compared to the PAPER Hex one. An important component of this plot is the averages over frequency in the right-hand panels. We saw that the average flagging for a given time sample was about 5% higher for HERA than for the PAPER Hex, mostly due to the higher occupancy of the FM band. But we also saw something new; HERA appeared to be much more sensitive to broadband bursts of RFI. The PAPER Hex caught one of these events (around 1.30am SAST) at high significance, but most of them hardly rose above average flagging. HERA saw five to seven bursts across the night.

6.2.2.4 Comparison PAPER-128 stacked flags

Section 6.2.1 presented RFI flags stacked over 150 days of observations. This method washed-out single events that effect analysis on a single-night basis, but was sensitive to repeatedly offending frequencies. Due to the PAPER-128 analysis pipeline, many channels were automatically flagged (particularly large portions of the band edges), which artificially boosted the average flagging per time and did not allow for closer inspection of the ends of the band. There was some evidence of broadband emission (see Figure 6.2) but the band was largely free of RFI in the middle of the night. Obviously, the data presented in this section shows a less-clean band, but it also only concentrated on a single




Figure 6.10: RFI flag waterfalls of frequency vs. South Africa Standard Time for the HERA Hex (*left*) and PAPER Hex (*right*). The top panels show the average over time (identical to Figure 6.9), while the right panels show the average over frequency.

night’s data, so it may be that IDR1 was conspicuous compared to an ‘average’ night of observations.

Traits shared between the two analyses were:

- Aircraft communications disrupting data until around local midnight.
- ORBCOMM spilling out of its band.
- VHF TV frequencies emitting throughout the night in the high end of the band.

Table 6.2.2.1 highlighted many frequencies seen by HERA and not by PAPER-128. Again, given the fact that flags were stacked and averaged in Section 6.2.1, these may not have been ‘new’, but they could have been. Particularly conspicuous were the emissions in the aeronautical radionavigation band.

6.2.2.5 Discussion

I have presented a first look at RFI in HERA-19 Commissioning data. Probably due to the height of the receiving element on HERA versus PAPER dipoles, much more RFI

was apparent, especially on the low and high ends of the band. Luckily, the EoR band was largely clean of RFI, except for an emitter at about 154 MHz, which could have corresponded to single-frequency mobile phone communications. Such communications are officially banned in the SKA Radio Quiet Zone (SKA Panel, 2012), which HERA is in the center of.

Only looking at a single night of RFI flags limited the predictive power of this study. More data will be required to establish whether or not this level of RFI was ‘normal’. Broadband RFI bursts require closer investigation.

Efforts to extend the HERA band to lower and possibly higher frequencies are currently under way. The FM radio band extends to around 65 MHz, while the VHF TV band extends to around 230 MHz, so the RFI environment should be a consideration for these efforts.

Meanwhile, I note that the RFI flagging routine used here, `xrfi_simple`, is indeed ‘simple’. More advanced RFI flagging algorithms such as AOFlagger (Offringa et al., 2012) should be tested in later studies.

6.3 Quality Assurance Metrics

Throughout reduction of PAPER-128 data, we found a number of ways for data to become corrupted, including failure of analog or digital components, incorrect cable connections, or improper feed installation. Many of these can cause failure of redundant calibration algorithms that are sensitive to non-redundancies within an array. In the case of a real-time calibration pipeline, such as the one implemented for HERA, it is essential to have quickly-generated metrics to assess the overall health of the array. Much of the heritage of PAPER-128 data processing is present in the HERA Real Time Calibration Pipeline (RTC; [Ali et al. \(in prep.\)](#)).

6.3.1 Mean Amplitude Flagging

The most critical and likely failure that an antenna could have was malfunctioning electronics losing power and temporarily “killing” it. This failure mode was characterized with unusually low signal coming from the antennas, causing the visibilities associated with that antenna to have much lower than average amplitude. This leads to the definition of the mean visibility metric for antenna i :

$$M_i = \frac{\sum_{j,v,t} |V_{ij}|}{N - 1} \quad (6.5)$$

where V_{ij} is the visibility for the baseline involving antennas i and j , N is the number of antennas in the array, and the sum is taken over all antennas j ($i \neq j$), times, and frequencies. When M_i is compared across the array, it can reveal antennas with anomalously low signal to noise.

Erroneously cross-polarized antennas were ones where the feed was rotated 90 degrees, or the cables were swapped for the two polarizations along the signal path (such a failure mode could be expected for a large array under active construction). This caused the linear polarization visibilities (EE or NN) to have a lower amplitude or correlation relative to the cross-polarizations (EN and NE). While these could appear to be dead antennas, the cross-polarization visibilities would show that one of the antennas in a given visibility is cross-polarized.

We defined the mean visibility cross-polarization metric as

$$P_i = \frac{M_i^{NE} + M_i^{EN}}{M_i^{NN} + M_i^{EE}}, \quad (6.6)$$

where M_i is defined in Equation 6.5 and are calculated for all the polarization pairs. If P_i is larger than some threshold, then antenna i may be cross-polarized. This method is effective when applied to long baselines, but for the shortest baselines in HERA, large-scale astrophysical polarization (e.g. Lenc et al., 2016) is observed by in all the instrumental visibilities (since, for example, the NN visibility is equal to the sum of the Stokes I and Q




Figure 6.11: An example of the mean amplitude flagging metrics on exercised on HERA-47 data. The top panel shows the mean amplitude flagging metric M_i , and the lower panel shows the cross-polarization metric P_i . The E feed of antenna 50 shows at $\sim 15\sigma$ deviation, causing it to be flagged from further processing.

visibilities). This brings P_i closer to unity for any antenna i .

Figure 6.11 shows the application of the metrics described above to raw HERA-47 data in the RTC system. Note that we plotted the modified z -score and not the raw metric. Traditionally the z -score is defined as the deviation from the mean divided by the standard deviation, and is a method to detect outliers. However, since our metrics were not necessarily normally-distributed, we use the modified z -score defined as

$$MZ_i = 0.6745 \frac{|x_i - \tilde{x}|}{MAD}, \quad (6.7)$$

where x_i is some data point (in our case this was a metric), \tilde{x} is the median of $\{x_i\}$, and MAD is the median absolute deviation. Iglewicz & Hoaglin (2010) recommend that modified z -scores with an absolute value greater than 3.5 should be flagged as outliers. In the analyses shown in Figure 6.11, we took the cut-off to be 5, to account for greater acceptable variation in our metrics.


The top panel of figure 6.11 shows the modified z -score for the mean visibility metric as a function of antenna number. Antenna 50E (East-West polarization) was a $\sim 15\sigma$ outlier, indicating a problem with amplification along the signal chain. Upon further inspection, antenna 50E was shown to exhibit strange behavior, dropping in and out of the signal chain. Further investigation is needed to know the root cause of this misbehavior.

The bottom row of figure 6.11 shows the modified z -score for the cross polarization metric as a function of antenna number. There were no modified z -scores outside of the [-5,5] range, indicating that there were no cross polarized antennas. Further inspection of the raw visibility data validated this finding.

6.3.2 Flagging on omnical χ^2

We developed a metric that identified days of PAPER-128 observations with poor overall omnical χ^2 values. Using the median of the χ^2 values in the EoR band (channels 100–160; 150–180 MHz), over all days of observation, we flagged days that exceeded the average of the median by one standard deviation from the mean. Figure 6.12 shows an example of this flagging method used on the latter half of the first PAPER-128 observing season. Clearly, there were just a few JDs that had exceptionally high χ^2 values. Our metric captures almost all of these outliers – the possible exceptions being JDs 2456717 and 2456718.

The omnical software also provides χ_a^2 values for each antenna a – that is, how much each antenna contributed to the overall non-redundancy of the array. The above sigma-clipping metric was applied to χ_a^2 , and was able to identify the same bad antennae using the mean amplitude flagging metric described in Equation 6.5.



Chapter 7

Polarimetric Calibration

Hypothetically, integrating over long observing seasons should average-down the noise in the EoR window, allowing for the study of the EoR without the risk of foreground contamination. However, the wedge-window paradigm arises due to the chromaticity of the interferometer as well as the signal. Spectral structure can be imparted to interferometric visibilities by the frequency evolution of the antenna beam and by calibration. Spectral structure can also be induced on otherwise-smooth foregrounds via Faraday rotation of polarized foregrounds. In of itself Faraday rotation is not a problem, since HI emission should be largely unpolarized, but errors in calibration and beam deconvolution can leak polarized signal into unpolarized visibilities (e.g. Thompson et al., 2017, ; Chapter 3).

Interferometers with N antennae that are sensitive to N_{pol} polarizations will generate $N_{\text{pol}}N(N - 1)/2$ visibilities per time-frequency sample, defined as:

$$V_{ij,pq}(\nu, t) = g_{p,i}^*(\nu, t)g_{q,j}(\nu, t)\exp(-2\pi\nu\tau_{pq}) \times \int d\Omega A_{pq}(\nu, \hat{s})S_P(\nu, \hat{s})\exp(i\vec{b}_{ij} \cdot \hat{s}\nu/c) \quad (7.1)$$

where $p, q \in x, y$ denote instrumental polarizations, i.e. the response to signal projected into the North-South direction or the East-West direction, or whichever direction the dipole arms of the instrument are oriented; i, j refers to two antennae with baseline

vector \vec{b}_{ij} , $S_P(v, \hat{s})$ is the sky temperature for Stokes parameter P , and $A_{pq}(nu, \hat{s})$ is the spatial sensitivity of the instrument to S_P (projected into the instrumental basis). Outside the integral are three direction-independent variables that must be calibrated: the complex gain of each dipole arm $g_{p,i}(v, t)$ and the phase between dipole arm p and q , τ_{pq} . For $p = q$, $\tau_{pq} = 0$.

All the above presents a data processing challenge: these visibilities must be precisely calibrated over long observing seasons to ensure that the cosmological signal is not averaged away by calibration errors and not contaminated by spectral structure. One way of overcoming part of this challenge is to construct large arrays of redundantly-spaced elements. The redundancy of the visibilities of such an interferometer allows the gain terms to be solved-for precisely by least-squares minimization algorithms (Liu et al., 2010). In Section 7.1, I explore the implications of implementing full-polarization redundant calibration on PAPER-128 data. However, redundant calibration is not the only tool radio astronomers posses that can be used to obtain precise calibration solutions. In Section 7.2, I present a basic implementation of full-polarization image-based calibration on data from the PAPER-32 polarized imaging array.

7.1 Redundant Calibration

Unpolarized redundant calibration has been pursued by the Donald C. Backer Precision Array for Probing the Epoch of Reionization (PAPER; Parsons et al. (2010)), the Hydrogen Epoch of Reionization Array (HERA; DeBoer et al. (2017)) and part of the Murchison Widefield Array (MWA; Tingay et al. (2013)). For EoR studies this has the advantage of being able to average-down noise in the EoR window once all redundant visibilities are calibrated, for potentially very high signal-to-noise measurements of narrow regions of k -space. However, it sacrifices uv -coverage, leading to poor imaging capabilities.

Redundant calibration of low-frequency interferometers has been demonstrated by Zheng et al. (2014) with the MIT EoR experiment, Parsons et al. (2014); Jacobs et al.

(2015); Ali et al. (2015) and Kolopanis et al. (2018) with PAPER and Li et al. 2017 with part of the MWA. All of these studies calibrated linearly-polarized instrumental visibilities.

Moore et al. (2017) used the same calibration parameters as Parsons et al. (2014) and Jacobs et al. (2015), but also solved for a single value of τ_{pq} for the observing season, since their analysis was for cross-polarized visibilities also. These PAPER studies took linear combinations of instrumental visibilities to form ‘pseudo-Stokes’ visibilities (e.g. Thompson et al. (2017), Moore et al. (2013)). In this Chapter, we use the notation and convention

$$\begin{pmatrix} V_I \\ V_Q \\ V_U \\ V_V \end{pmatrix} = \frac{1}{2} \begin{pmatrix} 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 1 & 0 \\ 0 & -i & i & 0 \end{pmatrix} \begin{pmatrix} V_{xx} \\ V_{xy} \\ V_{yx} \\ V_{yy} \end{pmatrix} \quad (7.2)$$

for these quantities.

As noted above, the spectrally-structured HI emission from the EoR should only be detected in V_I , but fractions of Faraday-rotated (spectrally structured) V_Q and V_U are capable of leaking into V_I via calibration errors and intrinsic properties of the complex beam. Short of detecting the polarized power spectrum of V_Q and V_U , which may be at levels beneath the EoR signal, we require confidence that we are accurately calibrating our data to prevent as much leakage as possible into V_I .

At these low frequencies and inherently large scales probed by modern instruments, the Stokes V sky appears to be empty. Indeed, Patil et al. (2017) used the spherical power spectra of their Stokes V images as a proxy for the thermal noise power spectrum. In wedge-space, one test of polarimetric calibration is therefore to see how close to thermal noise V_V is. Kohn et al. (2016) observed a direction-independent bias in V_V , but that study implemented a Moore et al. (2017)-style polarimetric calibration of τ_{pq} , solving for a single value across the array, which could indeed lead to a direction-independent bias (but was more powerful than not calibrating it at all; see their Figure 5).

In Section, we explore different schemes of redundant calibration which include polar-

ization. We present three different calibration schemes, all based around the OMNICAL¹ package (Zheng et al., 2014), and compare their uses and shortcomings. We structure the discussion as follows: in Section 7.1.1 we give a mathematical overview of the least-squares minimization algorithms implemented for the different calibration schemes and show basic simulations of each algorithm in action. We describe the real data used to test the calibration schemes, and the results of those tests, in Sections 7.1.2 and 7.1.3, respectively. We discuss our findings and conclude in Section ??.

7.1.1 Mathematical Overview

In this section we briefly describe the fundamental steps of the redundant calibration scheme implemented in OMNICAL. For a more thorough discussion of the algorithm, see Wieringa (1992), Liu et al. (2010), Zheng et al. (2014) and Dillon et al. (2017).

7.1.1.1 Redundant calibration with OMNICAL

We can express visibilities for baseline separation $|i - j|$, polarization pq as a system of equations, of form

$$V_{ij,pq} = g_{p,i}^* g_{q,j} V_{|i-j|,pq} + n_{ij} \quad (7.3)$$

where we have dropped the frequency and time dependence; this is true for every time-frequency sample. This system is overdetermined for a highly redundant array configuration. We assume that all polarizations have the same noise statistics for noise n_{ij} . An initial estimate for the least-squares fit can be obtained by solving the linearized equation in logarithmic space (termed *logcal*)

$$\log V_{ij,pq} = \log g_{p,i}^* + \log g_{q,j} + \log V_{|i-j|,pq} \quad (7.4)$$

which can constrain the parameter space, but produces biased results since noise is

¹<https://github.com/HERA-Team/omnical>

additive in linear space. Taylor expanding Equation 7.3 around the estimated values (denoted with bars) for each parameter as found by 7.4 grants a system of linearized equations (termed *lncal*)

$$V_{ij,pq} = \bar{g}_{p,i}^* \bar{g}_{q,j} \bar{V}_{|i-j|} + g_{p,i}^* \bar{g}_{q,j} \bar{V}_{|i-j|} + \bar{g}_{p,i}^* g_{q,j} \bar{V}_{|i-j|} + \bar{g} *_{p,i} \bar{g}_{q,j} V_{|i-j|} \quad (7.5)$$

which can be solved iteratively, minimizing the least-squares statistic w

$$w^2 = \sum_{\text{all } i \neq j} \sum_{p,q \in x,y} \frac{|V_{ij,pq} - g_{p,i}^* g_{q,j} V_{|i-j|,pq}|^2}{\sigma_{ij}}. \quad (7.6)$$

7.1.1.2 Including polarization in redundant calibration

We test three different calibration schemes for polarization.

- Use OMNICAL separately on the two linearly polarized visibilities V_{xx} and V_{yy} to obtain independent estimates for all gain terms, and apply those gains to all four instrumental polarizations of visibilities. We will refer to this scheme as *2pol* calibration.
- Include all four instrumental polarizations in the least-squares statistic. This means that, for example, information on the $g_{x,i}$ term will come from V_{xx} and V_{xy} . This will be referred to as *4pol* calibration.
- Use a scheme in which the cross-polarized visibilities V_{xy} and V_{yx} share values of the model visibility $V_{|i-j|,pq}$. This means that we are minimizing the value of $V_{yx} - V_{xy}$ during calibration; effectively minimizing the pseudo-Stokes V visibility. There is validity to assuming the pseudo-Stokes V visibility to be noise-like, as discussed above. We refer to this calibration scheme as *4pol+minV* calibration.

7.1.1.3 Degeneracies

One virtue of OMNICAL is that it can calibrate all antennas relative to one another without reference to an explicit sky model. However, this results in known degeneracies in the minimization of w^2 that cannot be resolved with OMNICAL and must be fixed afterwards with absolute calibration of the whole array using external information. The number of degeneracies is equal to the number of zero eigenvalues that arise in the linearized w^2 minimization algorithm. For a single polarization, there are 4 degenerate modes, which can be interpreted physically as overall amplitude, overall phase, phase tilt, and phase tip. Phase tilt and tip can be interpreted as a two-dimensional phase slope across the array. It is clear that these factors do not affect the product $g_{p,i}^*(\nu, t)g_{q,j}(\nu, t)V_{ij,pq}(\nu, t)$ (and we again refer the reader to Dillon et al. (2017) for a comprehensive description):

- Overall amplitude: $g_{p,i} \rightarrow Ag_{p,i}$ and $V_{ij,pq} \rightarrow V_{ij,pq}/A^2$.
- Overall phase: $g_{p,i} \rightarrow g_{p,i}e^{i\phi}$
- Two-dimensional phase slope: for a co-planar array, define a phase vector $\vec{\Upsilon} = (\Upsilon_X, \Upsilon_Y)$ where Υ_X, Υ_Y refer to Cartesian directions as opposed to polarizations. For an antenna at position \vec{v}_i , and defining $\vec{d}_{ij} = \vec{v}_i - \vec{v}_j$, then $g_{p,i} \rightarrow g_{p,i}e^{i\vec{\Upsilon} \cdot \vec{v}_i}$ and $V_{ij,pq} \rightarrow V_{ij,pq}e^{i\vec{\Upsilon} \cdot \vec{d}_{ij}}$. This is allowed because \vec{d}_{ij} is the same for all redundant visibilities.

OMNICAL ‘fixes’ the amount that the amplitude degeneracy is able to drift between samples by imposing that the average of the absolute of the gains over the array average to unity. Similar tricks can be played with the other degeneracies to project them into a space that does not adversely effect calibration.

The phenomenon of ‘degeneracy fixing’ in polarized redundant calibration was explored in depth by Dillon et al. (2017). In that work, they showed how the number of degeneracies changes with polarized calibration scheme, focussing on what we refer to as the *2pol* and *4pol* schemes in this work. We briefly review their results here.

In the *2pol* scheme, there are the 8 expected redundancies as expected for two independent calibrations; it is as if there are two co-located arrays for *xx* and *yy*, each with the four degeneracies listed above. However, in the *4pol* scheme, the number of degeneracies is reduced to 6. Introducing V_{xy} and V_{yx} breaks the phase tilt and tip degeneracies per polarization, and leaving a polarization-independent phase tilt and tip. Extending this formalism to the *4pol+minV* scheme², the number of degeneracies is further reduced to 5. By imposing equality between V_{xy} and V_{yx} , the ‘overall phase’ degeneracies per polarization are broken and there remains only a single overall phase for the array.

Dillon et al. (2017) emphasized the importance of understanding how fixing degeneracies can effect the amplitude of noise in redundant calibration solutions. They found that although the number of degeneracies decreased in the *4pol* scheme, V_{xy} and V_{yx} had low enough signal to noise that their inclusion in calibration greatly increased the noise in the gain solutions. Fixing only the 6 degeneracies of that scheme introduced much greater noise amplitude in the calibration solutions. However, fixing the 8 degeneracies from the *2pol* scheme *while still calibrating* with the *4pol* scheme gave almost identical gain solutions. Learning from this, we fix the 8 degeneracies from the *2pol* scheme throughout the rest of this work, independent of the calibration scheme used.

7.1.1.4 Expectations from simulation

Using the formalism of Nunhokee et al. (2017) and Chapter 3, and the HFSS³ complex voltage simulations of the PAPER beam described therein, we simulated the polarized response of the instrument to an unpolarized sky. That is, we passed a Stokes I - only sky model (de Oliveira-Costa et al., 2008) over a 30 m East-West baseline – the baseline vector of interest for the PAPER experiment (Parsons et al., 2014; Jacobs et al., 2015; Ali et al., 2015; Moore et al., 2017). This means that whatever was observed in the pseudo-Stokes Q, U and V visibilities could be interpreted as direction-dependent leakage from Stokes

²As explored in the public HERA Memo #30.

³<http://www.ansys.com/Products/Electronics/ANSYS-HFSS>

I (the expected regime for a single night of observations with PAPER; Kohn et al. 2016). The 2008 Global Sky Model used for the simulation was primarily diffuse in nature, with the ‘A-Team’ of bright low-frequency sources (Fornax A, Pictor A, etc.) also included. The diffuse nature was appropriate for the fringe profile of a 30 m baseline. The output of the simulation is shown in Figure 7.1 with and without the inclusion of an instrumental noise model, with the system noise drawn from Moore et al. (2017).

7.1.2 Data Processing

We tested these different calibration schemes on one night (JD 2456680.20 – .65; January 22nd-23rd 2014; 6pm – 6am South African Standard Time; Local Siderial Time 2 – 13.5 hours) of PAPER 128-element observations. The PAPER-128 signal chain and full observation season results will be discussed in forthcoming publications, but we will provide a brief overview here.

PAPER-128 consists of 128 dual-polarization dipole receivers, 112 of which are arranged in a highly-redundant configuration, and the rest placed as in- and outriggers to the array in order to improve *uv* coverage; see Figure 7.2. Since all dipole arms are oriented North-South (‘x’) and East-West (‘y’), *xy* and *yx* correlations have very low signal-to-noise compared to *xx* and *yy*.

All visibilities were RFI flagged using PYTHON scripts from the AIPY⁴ library. These took the derivative of the frequency axis of all baselines associated with a given antenna and flagged any frequencies with a derivative 6σ above the mean, per integration. We took the union of all baseline flags and applied them to the data. The 5 MHz on both band-edges were always flagged. Compression proceeded as described in Appendix A of Parsons et al. (2014), filtering to critical Nyquist sampling rates for the longest (300 m) baseline of 493 kHz along the frequency axis (203 channels) and 42.9 s along the time axis.

After RFI flagging and compression, *xx* and *yy* visibilities were checked for erroneous

⁴<https://github.com/HERA-Team/aipy>





Figure 7.1: Simulations of absolute value (*upper panels*) and phase (*lower panels*) of pseudo-Stokes visibilities (left to right: pseudo-Stokes I, Q, U and V), as measured by a 30 m East-West PAPER baseline. The upper group shows the noiseless simulation, and the lower shows the same simulation with the addition of a realistic PAPER noise model (Moore et al., 2017). Only a Stokes I sky was used – all of the structure seen in V_Q , V_U and V_V can be attributed to direction-dependent leakage.⁶³




Figure 7.2: Arrangement of the PAPER-128 array. Antennae determined to be malfunctioning are shown with red crosses and were excluded from analysis.

behavior including 2σ deviations from the median in number of RFI flags and mean visibility amplitude. Antennae exhibiting these behavior were excluded from further analysis (and are shown with red crosses in Figure 7.2). If an antenna qualified as ‘bad’ in one polarization, it was excluded in all of them. Also note that we can only calibrate the 112 antennae in the redundant grid using OMNICAL.

For the least-squares fit to converge at the *lbgcal* stage of calibration, visibilities cannot exhibit phase-wraps, since the system of equations solved at this stage are insensitive to additive offsets of 2π in their imaginary parts (*lincal* will be able to re-insert these as required; Liu et al. 2010). Therefore we must flatten the phases on all visibilities prior to any implementation of OMNICAL. We were able to do this redundantly without reference to the sky; taking the ratio of redundant (uncalibrated) visibilities together and averaging over time, defined as:

$$\mathcal{V}(\nu) = \langle V_{ij,pq}(\nu, t) V_{kl,pq}^*(\nu, t) \rangle_t. \quad (7.7)$$

Because this is the ratio of two nominally-redundant visibilities (i.e. baselines ij and kl belong to the same redundant group with model visibility $V_{|i-j|,pq}$), and the measurements are not yet calibrated, we can expand Equation 7.7 as

$$\mathcal{V}(\nu) = \langle g_{p,i}^* g_{q,j} g_{p,k} g_{q,l}^* |V_{|i-j|,pq}|^2 \rangle_t. \quad (7.8)$$

Writing the complex gains as $g_{p,a} = G_{p,a} e^{-i\nu\tau_{p,a}}$ for antenna a , we can reduce the product of gain-amplitudes and the squared visibility into some frequency-dependent function, and an exponential product of gain phase terms

$$\mathcal{V}(\nu) = K(\nu) \exp(i\nu(\tau_i - \tau_j - \tau_k + \tau_l)). \quad (7.9)$$

A Fourier transform along the frequency axis (which we term a *delay transform*; Parsons et al. 2012b) of $\mathcal{V}(\nu)$ gave a function that was sharply-peaked at a given delay

$$\tilde{\mathcal{V}}(\tau) = \tilde{K}(\tau) * \delta_D(\nu(\tau_i - \tau_j - \tau_k + \tau_l)). \quad (7.10)$$

We can define a variable as the maximum of the above function:

$$\mathcal{T}_{ijkl}(\tau) = \max |\tilde{\mathcal{V}}(\tau)|, \quad (7.11)$$

which will occur at value

$$\tau = \tau_{\max} = \nu(\tau_i - \tau_j - \tau_k + \tau_l). \quad (7.12)$$

With enough redundant baselines involving antennae $i, j, k \& l$ this is a linearly-solvable set of equations for each value of τ . Multiplying V_{ij} by $e^{-2\pi i\nu(\tau_i - \tau_j)}$ by definition flattened the phase across the band. This method was very sensitive to signal-to-noise, so these initial phase estimates were created with $p = q$; that is xx and yy visibilities only. The estimates were then applied to all visibilities appropriately. We could then run OMNICAL according to each of the schemes described in Section 7.1.1.2.

7.1.3 Results

We ran OMNICAL using the *2pol*, *4pol* and *4pol+minV* schemes, which granted complex gain values for each antenna feed in the redundant grid. In this work we chose to concentrate our analysis on the 30 m East-West spacings used for PAPER power spectrum studies.

7.1.3.1 Calibration

The complex-gains dataset alone was highly multidimensional. We chose to analyze short time- and frequency- averages for these data. For example, the $\langle |g_{x,a}^{2\text{pol}}| \rangle_{t,v}$ notation indicates the average of the absolute value of the gain value for antenna a , polarization ‘x’ in the *2pol* calibration scheme. The average is over 10 minutes (the length of a single MIRIAD file produced by the PAPER correlator) and a 10 MHz band running from 145 to 155 MHz (the center of the PAPER band, generally clear of RFI and used for power spectrum analyses).

Figure 7.3 shows $\langle \arg(g_a^{2\text{pol}}) \rangle_{t,v}$, the average phase of the gain calibration for ‘x’ and ‘y’ polarizations. It very clearly shows the phase-slope degeneracy present for both dipole orientations, sloping in opposite directions.

Figure 7.4 shows the differences in ‘x’ gain solutions between calibration schemes. The difference between *4pol* and *4pol+minV* was consistently smaller than the difference of either of these with the *2pol* scheme.

As noted in Section 7.1.1.3, OMNICAL tries to fix the average gain amplitude over the array to unity to avoid drifts in the amplitude degeneracy from sample to sample, so we can compare amplitude calibrations in terms of percentage deviation. The average difference in gain amplitude per antenna between *2pol* and *4pol* was 3.5%, between *2pol* and *4pol+minV* was 2.9%, and was 0.7% between *4pol* and *4pol+minV*. The ~ 30 degree spread in the differenced phases between the *2pol* scheme and the other two is likely due to different realizations of phase degeneracies between these calibration schemes.

Figure 7.5 shows the sum of w^2 values (see Equation 7.6) over all antennas in the




Figure 7.3: The phase of the complex gain value for the *4pol* calibration scheme is shown on the color axis (radians) for the redundant grid. The phase of the ‘x’ gains is shown on the left, and the ‘y’ gains on the right. The phase-slope degeneracy is clear in both panels, as is the fact that the slope is in opposite directions for the two dipole orientations.

array for each feed polarization in each of the three calibration schemes. The color scale is logarithmic. Clearly, the *2pol* scheme achieves a much greater level of redundancy in each feed polarization throughout the band. Towards the end of the night, the Galaxy is in the far side-lobes of the PAPER beam and introduces higher sky temperatures, which accounts for the trend in all calibration schemes performing worse towards the end of the night. However, the reason for rapid transitions in w^2 at Local Sidereal Time ~ 10.5 in the *4pol* and *4pol+minV* schemes is not well understood. In general, w^2 is an order of magnitude higher for the *4pol* and *4pol+minV* schemes.

7.2 Imaging Calibration




Figure 7.4: Histograms of differences in the ‘x’ gain calibrations per antenna between calibration schemes. Above: difference in absolute value. Below: difference in phase. For most antennae, absolute value of the gain does not change by large amounts between calibration schemes; most of the change takes place in the phase. The difference between *4pol* and *4pol+minV* is consistently smaller than the difference of either of these with the *2pol* scheme. This is likely a sign of different realizations of phase degeneracies between calibration schemes.




Figure 7.5: w^2 values, summed across the array, for the x (above) and y (below) feeds in different calibration schemes. From left to right: *2pol*, *4pol* and *4pol+minV*. The color axis is logarithmic, in arbitrary data units. The white gaps are due to RFI flagging.

Chapter 8

The Ionosphere

The ionosphere is a section of Earth's atmosphere composed of several layers, between 60 and 1000 km in altitude. It overlaps the Troposphere, Stratosphere, Mesosphere, Thermosphere and Exosphere. The ionosphere is an ionized plasma, composed of ions from molecules in the atmospheric layers it overlaps that are ionized by solar radiation. The ionization state of the ionosphere can be quantified by the Total Electron Content (TEC) – an integral of electron count in a given direction – among other metrics.

Spatiotemporal variations of the TEC are tied to solar activity, and therefore largely both diurnal and seasonal. More ionization, and therefore a larger TEC, is to be expected in the day time and closer to the summer solstice. The Solar Cycle also influences TEC, with more sunspots proportional with a higher TEC; at solar maximum, this effect dominates the seasonal variation (Sotomayor-Beltran et al., 2013). Ionospheric variations are typically described as Kolmogorov turbulence (i.e. small scale motions are isotropic in their direction and scale with wavenumber; Zolesi & Cander 2014), however, LOFAR observations report deviations from isotropy in their observations (Intema et al., 2009; Mevius et al., 2016). Regions of the ionosphere that can be assumed to be constant in density and shape at a given time are referred to as "isoplanatic patches". At 74 MHz, these patches are observed to be $1^\circ - 2^\circ$ in radius (Cotton & Condon, 2002).

The ionosphere is composed of three main layers: D, E and F, which vary according to

Table 8.1. Ionospheric Layers

Layer	Time	Altitude km	Components	Electron Density $e^- m^{-3}$
D	Day	60–90	NO^+ , N_2 , Ar , O_2^-	$10^8 - 10^9$
E	Day/Night	90–150	NO^+ , O_2^+ , O^+ , N_2^+	10^{11}
F_1	Day	140–600	NO^+ , O_2^+ , O^+ , N^+	10^{11}
F_2	Day/Night	220–800	O^+ , H^+ , He^+	$10^{10} - 10^{13}$

the day-night cycle. These are summarized in Section 8 (which summarizes Chapter 2 of Zolesi & Cander 2014). At night, there are not enough high-energy electrons to penetrate to lower altitudes, causing the D layer to recombine. The E layer increases in altitude at night due to a similar effect. The E and F layers persist at all times, but during daylight the F layer is divided into two sub-layers, F_1 and F_2 .

The diurnal nature of the ionosphere is important to radio propagation. During the day, the D layer reflects radio transmissions much closer to the Earth than during the night, when the E and F layers reflect. This leads to longer-range transmissions being possible after sunset¹.

The relevance of the ionosphere to this work is its coupling with Earth's magnetic field. Recall that, as mentioned in previous chapters, a linearly polarized electromagnetic wave, propagating through an ionized plasma which has an incident magnetic field, will experience Faraday Rotation of its original polarization angle χ :

$$\chi_{\text{obs}} = \chi + \phi \lambda^2 \quad (8.1)$$

where λ is the wavelength, and

$$\phi(\hat{s}) \approx 0.81 \int_{\text{source}}^{\text{obs}} n_e(\hat{s}) \vec{B}(\hat{s}) \cdot d\vec{s} \quad (8.2)$$

¹This effect was first observed by E. V. Appleton (Appleton, 1946), confirming the ionosphere's existence, for which he was awarded the 1947 Nobel Prize in Physics.

where the source of the electromagnetic wave is in direction \hat{s} on the sphere, n_e is the electron density scalar field and \vec{B} is the magnetic vector field. The Rotation Measure (RM) ϕ is the integral of the product along the line of sight, and has units of rad m⁻². Since the ionosphere is capable of imparting an additional RM to polarized radio waves, inducing spectral structure to interferometric visibilities, understanding it is crucial to quantifying the effect of polarization on EoR measurements.

In this chapter, I review historical measurements of the ionospheric TEC and RM distributions in Section 8.1 and modern observations in Section 8.2. In Section 8.3 I present our work on the role of the ionosphere in PAPER and HERA measurements, and software we developed to quantify those effects.

8.1 Historical measurements of TEC and RMs

The existence and layered nature of the ionosphere was confirmed between the 1920s and the 1940s. Measurements of the TEC and RM distributions came later, once radio-communications satellites were put in orbit, and are closely tied to the Global Positioning System (GPS) launched in the late 1970s (called the NAVSTAR system). NAVSTAR GPS satellites transmit at two narrow frequency bands, centered about 1.2276 GHz (“L₂”) and 1.57542 GHz (“L₁”). Encoded in these transmissions are the local clock times per satellite (precisely calibrated with one another and with ground clocks) and their positions. With four satellites in view of a receiver, one is capable of computing their three-dimensional position and their local clock relative deviation from the satellite clock time.

MacDoran & Spitzmesser (1989) showed that one could use a frequency-dependent time delay induced by the ionospheric plasma (Klobuchar, 1983; Brunner & Welsch, 1993):

$$\Delta t_{\text{iono}} = \frac{40.3}{cv^2} \text{TEC} \quad (8.3)$$

to calculate an estimate of the TEC in the direction of a GPS satellite. Their approach

has been continuously refined. Using an estimate of the polarization angle of the emitted L_{1,2} transmissions, Titheridge (1972) and Royden et al. (1984) presented measurements of TEC by measuring the Faraday Rotation induced and worked towards an estimate of the TEC based on the RM. Lanyi & Roth (1988) showed that the more accurate method was calculation of the TEC using Δt_{iono} from Equation 8.3. Mannucci et al. (1998) introduced the Ionosphere Map Exchange Format (IONEX): a method and file format for storing TEC measurements using GPS beacons across the globe, allowing the first global TEC maps to be calculated. IONEX files contain global TEC measurements with a 2 hour cadence and generally 5° by 2.5° resolution in longitude and latitude respectively. They neglect the layered nature the ionosphere, modelling it as a thin sheet. Iijima et al. (1999) provided a server that automatically pushed IONEX files to the World Wide Web as soon as they could be constructed. Komjathy et al. (2005) presented the first measurements with over 1000 GPS stations. Recently, Erdogan et al. (2016) presented a method for time-series forward modelling of the TEC distribution using IONEX files.

Meanwhile, many generations of the International Geomagnetic Reference Field (IGRF ?) have continually improved the model of the Earth's magnetic field. This model is composed by spatial interpolation of magnetic field measurements (in up to 13th-order spherical harmonic coefficients) reported by institutions around the world.

Combining these two measurements – IONEX and IGRF data – can provide a map of RM distribution above any given position on Earth to moderate precision (better in the Northern Hemisphere than the Southern one, based on the number of GPS beacons in each). Afraimovich et al. (2008) offered the first such software implementation, with the objective of using it to track Solar Activity². Sotomayor-Beltran et al. (2013) introduced the `ionFR` package, which calculated ionospheric RMs towards a given position on the sky. We generalized their approach for the wide-field measurements in our `radionopy` software package, which we present in Section 8.3.

²Erickson et al. (2001) were the first to present software capable of calculating ionospheric RMs using the IGRF, but they used local GPS beacons instead of IONEX files

8.2 Low frequency observations: discoveries and challenges

Low frequency interferometric observations are effected in two main ways by ionospheric turbulence: scintillation in Stokes I observations, and Faraday Rotation in Stokes Q and U observations.

TEC variations introduce a variable index of refraction across a field of view. Stokes I signal from a point source will scintillate, change position, by an amount (e.g. Thompson et al., 2017):

$$\Delta\theta = -\frac{1}{8\pi^2} \frac{e^2}{\varepsilon_0 m_e} \frac{1}{v^2} \nabla_{\perp}(\text{TEC}) \quad (8.4)$$

at the observed frequency v , where ∇_{\perp} is the transverse gradient in TEC towards the direction of the source. The time, space and frequency dependence of this effect causes difficulty for long integrations, since the scintillation will cause averaging of point sources with empty space, spreading-out their signal over a $\sim \Delta^2\theta$ area. This can be interpreted as an additional source of noise in a Stokes I map. Vedantham & Koopmans (2015) showed that this scintillation noise can be much larger than image noise for baselines longer than ~ 200 m. Vedantham & Koopmans (2016), extending the previous analysis to the Fourier domain, showed that this noise does not pose large issues to HERA or SKA-Low EoR efforts, since realistic amounts scintillation were not sufficient to wash-out EoR signals on large scales (their dense cores of relatively short baselines also help). However, it could pose large issues for point-source calibration and subtraction methods – as emphasized in a public SKA memo by Cornwell (2016).

Loi et al. (2015) used MWA observation snapshots to map the scintillation as a function of space and time, resulting in the discovery of “tubes” of plasma density waves across the Southern Hemisphere in lines of roughly constant latitude. Comparing the sources in their snapshot images to source positions in the NRAO VLA Sky Survey (NVSS Condon et al., 1998) they were able to calculate displacement vectors, and showed

that they were strongly aligned to Earth’s magnetic field.

The literature surrounding ionospheric Faraday Rotation is less extensive than work focussing on the unpolarized component. Lenc et al. (2016) showed that MWA measurements of diffuse foregrounds could provide a map of ionospheric spatiotemporal variance as their RM changed throughout a series of observations. Lenc et al. (2017) showed that point source power could be seen “twinkling” in and out of polarized intensity maps due to ionospheric activity.

8.3 Relevance for PAPER and HERA EoR measurements

Within the PAPER and HERA power spectrum pipelines, many tens to hundreds of days of visibilities are averaged over during binning in LST. The ionosphere-induced spatial and temporal fluctuations in RM could produce sufficient phase scrambling of the celestial Faraday-rotated, polarized signal to suppress a fraction of any polarized signal leaked by some mechanism into Stokes I measurements. The fringe size of the 30 m baselines used in power spectrum analyses is large enough that scintillation effects are negligible.

This effect was first investigated in Moore et al. (2017). Using the `ionFR` package (Sotomayor-Beltran et al., 2013) we calculated the RM distribution at a single zenithal pointing throughout the PAPER-32 observation season. This was a vast simplification given the PAPER primary beam was much larger than a typical isoplanatic patch. Shown in Figure 8.1, there was a large spread of ionospheric RMs for each LST. There was a decrease in the average magnitude of the RM as LST increased. This was expected, given the strong correlation between the day/night cycle and TEC values (e.g. Tariku, 2015), and given that for this observing season, LST=4 hr corresponded to observations taken shortly after sunset, while LST=8 hr was always well into the night.

Treating the single pointing as constant over the sky, we calculated the expected attenuation of polarized signal, leaked into pseudo-Stokes I visibilities, that would be averaged over varying ionospheric conditions during LST binning. These attenuation factors were




Figure 8.1: Distribution of zenithal ionospheric RMs for 3 LSTs in the PAPER-32 observing season. From top to bottom: a histogram of the zenith ionospheric RMs over the season, for the transit of LSTs 4, 6, and 8 hr. Taken from Moore et al. (2017).

$43 \pm 6\%$ at 165 MHz and $7 \pm 5\%$ at 126 MHz.

To build on this result, we required more sophisticated simulations of the interaction of the polarized sky with the instrument and whole-sky maps of the ionosphere. To accomplish the latter, we developed the open-source Python package `radionopy`³. Like `ionFR`, `radionopy` uses GPS-derived TEC maps from IONEX files and the IGRF to estimate the value of ionospheric RM at a given latitude, longitude and date. Unlike its predecessor, `radionopy` does not necessarily calculate an RM at a given pointing, but instead is capable of calculating the ionospheric RM over a HEALPix grid of the sky (?). Such an expression of ionospheric variation is natural to wide-field, drift-scanning EoR arrays, and reflects the format of the IONEX input measurements, which are given in their spherical harmonic decompositions. `radionopy` is vectorized, leading to efficient generation of full-sky ionospheric maps, and object-oriented, allowing for easier collaborative development. Additionally we implemented the interpolation scheme recommended in the IONEX documentation to obtain “best-guess” full-sky maps for arbitrary times between the 2-hour time resolution of IONEX data.

An example output from `radionopy` is shown in Figure 8.2 as a HEALPIX grid of the hemisphere observable from the PAPER site in the Karoo. In Figure 8.3 we show `radionopy` and `ionFR` output for a single pointing towards Cassiopeia A (Cas A; RA= $23^{\text{h}}23^{\text{m}}27.9^{\text{s}}$, Dec= $+58^{\circ}48'42.4''$) from the LOFAR Core site in the Netherlands. The two codes gave qualitative agreement. Slight offsets at the highest RM values that day could be attributed to differences in our interpolation schemes.

[Martinot et al. \(in prep.\)](#) investigated the full interaction of the polarized sky with the ionosphere, using realistic polarized sky models and fully-polarized HERA beam models (see Chapter 10 for an example). Their work revealed that the Moore et al. (2017) analysis overestimated the ionospheric attenuation due to their single-pointing and simple beam models. Realistic levels of attenuation for a 100 day HERA integration can be expected to reach a factor of ≤ 0.1 . If polarization leakage occurs close to the EoR level, this is

³<https://github.com/UPennEoR/radionopy>




Figure 8.2: An example of widefield ionospheric RMs calculated by radionopy.

sufficient to recover the EoR power spectrum. However, if it is above the EoR level (as expected by Nunhokee et al. 2017), the ionosphere alone will not be sufficient to rule out polarization leakage being detected before the EoR can be recovered.




Figure 8.3: The RM of Cas A as viewed from the LOFAR Core site in the Netherlands on April 11th, 2011, according to `ionFR` and `radionopy`. The two codes show quantitative agreement, demonstrating that `radionopy` can be used for single-pointing as well as full-sky RM measurements.

Chapter 9

A view of the EoR window from the PAPER-32 imaging array

In this Section, we present 2D power spectra created from data taken by the PAPER-32 imaging array in Stokes I, Q, U and V.

The PAPER 32-antenna array relied on its highly redundant configuration in order to take the measurements resulting in the strong upper limits on the 21 cm power spectrum (Parsons et al., 2014; Jacobs et al., 2015; Moore et al., 2017). However, for three nights in 2011 September, the 32 elements were reconfigured into an polarized imaging configuration. Creating power spectra allowed us to observe and diagnose systematic effects in our calibration at high signal-to-noise within the Fourier space most relevant to EoR experiments. We observed well-defined windows in the Stokes visibilities, with Stokes Q, U and V power spectra sharing a similar wedge shape to that seen in Stokes I. With modest polarization calibration, we saw no evidence that polarization calibration errors moved power outside the wedge in any Stokes visibility, to the noise levels attained. Deeper integrations will be required to confirm that this behavior persists to the depth required for EoR detection.

The layout of this Chapter is as follows. In Section 9.1 we provide a brief description of the PAPER array in its imaging configuration, the data from which this paper is based,

and describe its calibration and reduction. We also describe the method used to create 2D power spectra in this section. We analyze the power spectra in Section 9.2, and discuss the implications of our findings and conclude in Section 9.3.

9.1 Observations & Reduction

We present measurements taken overnight on 2011 September 14–15 over local sidereal times (LSTs) 0–5 hr.

Antennae were arranged in a pseudo-random scatter within a 300 m-diameter circle, the layout of which is shown in Figure 9.1. This allowed us to obtain resolutions between 15' and 25' across the bandwidth (100–200 MHz nominally, although in reality this extends 110–185 MHz due to band edge effects and VHF TV). Drift-scan visibilities were measured every 10.7 s, and divided into datasets about 10 minutes in length. We express an interferometric visibility V_{ij}^{pq} between antennae i (with dipole arm p , which can be x (East-West) or y (North-South) for PAPER dipoles), and j (with dipole arm q), in directional cosines l and m for frequency ν at time t , as:

$$V_{ij}^{pq}(\nu, t) = g_i^p g_j^{q*} \exp(-2\pi i \nu \tau_{pq}) \times \int d\Omega A^{pq}(\Omega, \nu) S(\Omega, \nu) \exp\left(\frac{-i\nu}{c} \vec{b}(t) \cdot \hat{s}(\Omega)\right) \quad (9.1)$$

where the g terms represent the complex gains for each antenna and dipole arm, A^{pq} is the polarized beam and S is the sky. The product $\vec{b}(t) \cdot \hat{s}(\Omega)$ represents the projection of the baseline between i and j with respect to an arbitrary location on the sky. The motivation for including the term for the delay between dipole arms p and q , τ_{pq} , is given in Section 9.1.1.3. This delay is clearly zero if $p = q$.

Visibilities were obtained from correlating both x and y dipoles, forming V^{xx} , V^{xy} , V^{yx} and V^{yy} . Frequencies from 100 to 200 MHz were sampled into 2048 channels. Data were delay-filtered to 203 frequency channels (see the Appendix of Parsons et al., 2014) and Chapter 6. Cross-talk was modelled and removed by subtracting the average power over




Figure 9.1: The PAPER-32, dual-pol antenna imaging configuration (top). They were arranged in a pseudo-random scatter within in a ~ 300 m diameter circle to maximize instantaneous uv coverage (bottom). uv coverage is shown for 100–200 MHz over 203 channels in blue, and 146–166 MHz over 20 channels in red (the latter being the frequencies used in our power spectrum analysis). Malfunctioning antennae identified during calibration are overlaid with red crosses (and are excluded from the uv coverage map).

the 5 hours of observation, which extended across LST=0h–5h. An initial RFI-flagging removed any outliers more than 6σ from a spectrally smooth profile.

9.1.1 Calibration

Calibration took place in three stages, detailed below: a first-order delay-space calibration for the initial gains and phases with respect to Pictor A, an absolute calibration using imaging with respect to Pictor A and Fornax A, and a polarimetric correction for the τ_{xy} phase term in the V^{xy} and V^{yx} visibilities. Traditional polarimetric calibration proceeds by observing a source with a known polarization angle, and solving for up to seven direction-independent terms in the Jones matrix (e.g. Thompson et al., 2017; Hamaker et al., 1996), as well correcting for the effects of the primary beam. Given the dearth of suitable calibrators at our observing frequencies, especially at the relatively low resolution and sensitivity of the array, we proceeded with polarized calibration using different techniques, as described in Section 9.1.1.3 below.

9.1.1.1 Initial calibration

A first-order gain and phase calibration was performed by a similar approach to Jacobs et al. (2013). Each 10 minute drift-scan dataset was phased to the known position of Pictor A using aipy routines.

The gain term in Equation 9.1 was approximated as

$$g_i^p = G_i^p \exp(-2\pi i v \tau_{ip}) \quad (9.2)$$

and the required delay τ_{ip} offset of the uncalibrated delay tracks to the real position on the sky solved for to obtain a phase calibration; the absolute flux calibration G_i^p was found by isolating the tracks of Pictor A in delay space, and applying the required flux scale across the band (for a discussion of delay-space calibration, see Parsons et al., 2012b, and Figure 9.4).

9.1.1.2 Absolute Calibration

Visibilities were converted to CASA Measurement Sets to be further calibrated using a custom pipeline developed around CASA libraries. Snapshot images were generated for

each 10 minute observation by Fourier transforming the visibilities. We used uniform weights and the multi-frequency synthesis algorithm to further improve the *uv* coverage. Dirty images were deconvolved down to a 5 Jy threshold using the Cotton-Schwab algorithm. The sky model generated by the CLEAN components was used to self-calibrate each snapshot over the full bandwidth, using a frequency-independent sky-model and averaging over the 10 minute observation. We corrected for residual cable length errors by computing antenna-based phase solutions for each frequency channel for each snapshot observation. After self-calibration, snapshot visibilities were again Fourier transformed into images and deconvolved down to a 2 Jy threshold to form the final sky models. These final sky models were used to solve for a frequency independent, diagonal, complex Jones matrix (Hamaker et al., 1996; Smirnov, 2011) for each antenna in order to calibrate gain variations from snapshot to snapshot. We make no attempt to correct sky models for polarized primary beams and, therefore, our gain solutions incorporate both the direction independent and the direction dependent responses of the two gain polarizations. This is a reasonable approximation for the scope of the paper, as, eventually, wide-field polarization corrections cannot be implemented directly in the per-baseline power spectrum estimation (see Section 9.2).

The average correction in magnitude through this second-order calibration was a $\pm 6\%$ change for *x* gains and $\pm 7\%$ for *y* gains from those derived in the initial delay-space calibration. If the gain on an antenna deviated by more than 30% from image-to-image during this analysis, it was discarded from future processing stages, since it was likely malfunctioning. This was true for 3 antennae (see the top panel of Figure 9.1).

The final gain amplitude calibration was carried out similarly to Ali et al. (2015). We generated single channel images between 120 and 174 MHz for each snapshot and deconvolved each of them down to 10 Jy. For each snapshot, a source spectrum is derived for Pictor A by fitting a two dimensional Gaussian the source using the PyBDSM¹ source extractor (Mohan & Rafferty, 2015). Spectra were optimally averaged together

¹<http://www.lofar.org/wiki/doku.php?id=public:usersoftware:pybdsm>

by weighting them with the primary beam model evaluated in the direction of Pictor A. To fit the absolute calibration, we divided the model spectrum (Jacobs et al., 2013) by the measured one and fit a 6th order polynomial over the 120-174 MHz frequency range. This procedure was repeated using Fornax A with the only difference that a taper was applied to the visibilities (120 m) in order to reduce Fornax A to a point-like source and use the model spectrum from Bernardi et al. (2013). The best fit coefficients for Pictor A and Fornax A were averaged together to obtain the final absolute flux density calibration. Snapshots of fully CASA-calibrated data are shown in Figure 9.2.

9.1.1.3 Polarimetric factors

Standard full polarization calibration involves correcting for leakage of Stokes I into the V_{ij}^{xy} and V_{ij}^{yx} visibilities and leakage of polarized signal into total intensity (the so called Jones D matrices or D -terms; e.g. Thompson et al. (2017); Hamaker et al. (1996)), and an unknown phase difference between the x and y feeds (e.g. Sault et al., 1996).

We attempt no D matrix calibration in this paper, as there is not a dominant source to be used for such calibration: the limited sensitivity of our observations does not offer good signal-to-noise ratio on PMN J0351-2744, the only polarized source at low frequencies known so far in our survey area. In addition, D -term calibration would require determination of the primary beam Mueller matrices beyond our current accuracy. The consequences of this limitation are discussed in the analysis of our power spectra in Section 9.2.

As an intermediate measure compatible with these limitations, we therefore adopted a minimization of the phase difference between the V_{ij}^{xy} and V_{ij}^{yx} visibilities, minimizing a sum of squared weighted residuals w :

$$w(v, t, \tau_{xy}) = \sum_{ij} |V_{ij}^{xy} - V_{ij}^{yx} \exp(-2\pi i v \tau_{xy})|^2 \quad (9.3)$$

to find an estimated value of τ_{xy} for the array at each (v, t) sample. This is equivalent to assuming that the sky is intrinsically not circularly polarized at the frequencies observed

by PAPER.

We choose not to correct for ionospheric Faraday rotation in our calibration. Not only is this difficult to do for widefield instruments, but also the ionosphere was relatively stable during the observations, so we expect little incoherent averaging during the power spectrum stage below. We calculated the stability of ionospheric RM (ϕ_{iono}) using the IONFR software (Sotomayor-Beltran et al., 2013), which calculates the ϕ_{iono} for a given longitude, latitude and time by interpolating values of GPS-derived total electron content maps and the International Geomagnetic Reference Field (Finlay et al., 2010). The values of ϕ_{iono} for different lines of sight are shown in Figure 9.3. Fluctuations of ϕ_{iono} will cause incoherent time-averaging and subsequent loss of polarized signal. Using the formalism of Moore et al. (2017) to calculate the attenuation factor, we found that none of the lines of sight (except for the 21h,0° one which goes beneath the horizon) shown are responsible for attenuating signal by > 20% in power-spectrum space (see Section 9.1.2).

We form linear combinations of the instrument visibilities, the so-called pseudo-Stokes visibilities (see e.g. Moore et al., 2013) V^I , V^Q , V^U and V^V as:

$$\begin{pmatrix} V^I \\ V^Q \\ V^U \\ V^V \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 1 & 0 \\ 0 & -i & i & 0 \end{pmatrix} \begin{pmatrix} V^{xx} \\ V^{xy} \\ V^{yx} \\ V^{yy} \end{pmatrix} \quad (9.4)$$

Data that were reduced, calibrated, and formed into Stokes visibilities were separated into delay spectra inside and outside of the horizon for each baseline. We used a 50 ns margin for what was considered ‘inside’ the horizon, in order to confine all supra-horizon emission (e.g. Parsons et al., 2012b; Pober et al., 2013) to the foreground component of the data. We implemented a one-dimensional CLEAN (Parsons & Backer, 2009; Parsons et al., 2012a) with a Blackman-Harris window to a tolerance of 10^{-9} . RFI is more easily identified in foreground-removed data, so we RFI-flagged again on the background data deviations greater than 3σ . We then added the inside- and outside-horizon visibilities back together; RFI flags were preserved in the process.

The effect of our calibration is shown in the delay-transformed visibilities in Fig-

ure 9.4. As is apparent in Figure 9.2, after improved calibration there are fewer delay tracks (i.e. sources) in the Stokes Q visibilities, while there is little overall change in Stokes U. The minimization of Stokes V, performed after the imaging calibration stage, moves power from Stokes V into Stokes U, effectively accounting for part of a D -term correction. But without an accurate D -term calibrator, Stokes U exhibits additional (and dominant) D -term leakage from Stokes I, in this case due to Pictor A. Pictor A is the brightest source in Stokes I in our observed field, and thus dominates the visibility shown. There is no reason to suppose that Pictor A is pure Stokes U (compare also Figure 9.2), and thus the bulk of this emission must be leakage.

9.1.2 Creating power spectra

Expressing the visibility $V_{ij}^{pq}(v, t)$ observed at time t (see Equation 9.1) in terms of the geometrical delay $\tau_g = \vec{b}(t) \cdot \hat{s}(l, m)/c$ for the baseline ij , Parsons et al. (2012b) define the delay transform as the Fourier transform of the visibility along the frequency axis:

$$\tilde{V}_{ij}^{pq}(\tau, t) = \int dv V_{ij}^{pq}(v, t) e^{2\pi i v \tau} \quad (9.5)$$

We can represent the power at each frequency and baseline in an array as a power spectrum in terms of their respective Fourier components k_{\parallel} and k_{\perp} as:

$$P(k_{\parallel}, k_{\perp}) \approx |\tilde{V}_{ij}^{pq}(\tau, t)|^2 \frac{X^2 Y}{\Omega B} \left(\frac{c^2}{2k_B v^2} \right)^2 \quad (9.6)$$

where B is the bandwidth, Ω is the angular area (i.e. proportional to the beam area), and X and Y are redshift-dependent scalars calculated in Parsons et al. (2012a).

To form $|\tilde{V}_{ij}^{pq}(\tau, t)|^2$, consecutive integrations were cross-multiplied, phasing the zenith of latter to the former i.e.:

$$|\tilde{V}_{ij}^{pq}(\tau, t)|^2 \approx |V_{ij}^{pq}(\tau, t) \times V_{ij}^{pq}(\tau, t + \Delta t) e^{i\theta_{ij,zen}(\Delta t)}|^2 \quad (9.7)$$

where $\Delta t = 10.7$ seconds and $\theta_{ij,zen}(\Delta t)$ is the appropriate zenith rephasing factor. This method should avoid noise-biased power spectra except on very long baselines, which

the PAPER configuration does not contain, while sampling essentially identical k -modes. Note that this is the same method used by Pober et al. (2013) in their investigation of the unpolarized wedge.

9.2 Results

Combining visibilities using Equation ??, we formed power spectra over frequencies 146–166 MHz according to Equation 9.6 using consecutive integrations for each Stokes visibility over time, and gridded our results into k -space, averaging in time. The k_{\perp} -axis is binned with a resolution of $4.65 \times 10^{-4} h \text{ Mpc}^{-1}$ to slightly reduce gaps in k -space due to missing baselines. This gave an average bin occupancy of 1.7 ± 0.9 . The resolution in k_{\parallel} ($5.06 \times 10^{-4} h \text{ Mpc}^{-1}$) is set by the 20 MHz bandwidth with 500 kHz resolution that we use in this analysis. Note that the Blackman-Harris window used in the delay-filtering stage after forming Stokes visibilities correlates adjacent frequency bins, and hence k_{\parallel} bins. Each $(k_{\perp}, k_{\parallel})$ bin was normalized by its occupancy.

Two-dimensional power spectra have been proven as powerful tools for large dipole array experiments, not only for assessing cosmology but also in order to constrain instrumental and analytical systematics (e.g. Morales et al., 2012). Polarization axes are a useful addition for such analyses, since we expect Stokes I to be approximately 3 orders of magnitude stronger than the other polarization products at the low radio frequencies and tens-of-arcminute scales native to PAPER observations (e.g. Pen et al., 2009; Moore et al., 2013) and when observing far from the Galactic Plane. This alone allows us to assume that much of the structure in the power spectra with power comparable to Stokes I is leakage. As we explore below, these leakage terms can come from direction-dependent effects (e.g. wide-field beam leakage; Carozzi & Woan, 2009) or direction independent ones (e.g. Mueller matrix mixing via gain errors and D -terms; Thompson et al., 2017) and appear with high signal-to-noise in power spectra.

Figure 9.5 shows power spectra in ‘pitchfork’ form (Thyagarajan et al., 2015b,a), with

k_{\parallel} in negative and positive directions (according to the East and West horizons, marked in white (horizon) and orange (horizon+50 ns delay, respectively). Each Stokes parameter pitchfork has its own interesting characteristics, which allow us to analyze different sky and instrument behaviors. The ‘wedges’ described in the literature that define the EoR window are simply the average of negative and positive values of k_{\parallel} . While we focus on the pitchfork expression of the power spectra in our results, we also show them in wedge form in Figure 9.6².

Simplifying the results of Thyagarajan et al. (2015b,a, see their papers for a full discussion), we expect power from diffuse emission to appear at low values of k_{\perp} and high values of k_{\parallel} , while point sources lie at all k_{\parallel} (all over the sky) but are down-weighted by the primary beam, which is broad, leaving a concentration of the power close to the $k_{\parallel} = 0$ line.

In Stokes I, we see the strongest power on most baselines arising at values $k_{\parallel} \approx 0$. This is expected in a situation of point sources that are relatively bright compared to any diffuse emission. Indeed, at the LSTs we observed at, several unresolved bright point sources transit the field (e.g. Figure 9.2), while the dominant source of diffuse emission at these frequencies, the Galactic plane, was below the horizon. However, we do see strong super-horizon emission at $0.02 \leq k_{\perp} \leq 0.03$, biased towards negative k_{\parallel} values. There is also a decrease in power with increasing k_{\perp} – both of these effects are consistent with the Thyagarajan et al. (2015a) simulations of faint diffuse structure transiting zenith.

The Stokes Q wedge shows a concentration of power close to $k_{\parallel} \approx 0$, similar to Stokes I. The inherent low polarization fraction at our frequencies works in our favor in detecting gain errors, since Stokes Q is largely expected to be faint, and thus the gain errors causing leakage from I appear at high signal-to-noise there. Indeed, this power decreases noticeably with more accurate gain amplitude calibration, but bright streaks at specific values

²Note the difference in the power distribution within the horizon differs from that shown in the Pober et al. (2013) V^{yy} wedge. That study used the PAPER 64-element, single-polarization imaging array to create power spectra in a ‘loud’ field containing point sources and Galactic signal, causing their wedge to be ‘fuller’ than the ones presented in this study.

of k_{\perp} remain, suggesting lower-level residual gain calibration errors on select baselines. Another possible source of power in Stokes Q stems from wide-field direction-dependent gain errors causing a non-smooth evolution of the sources on the edges of the beam. However, we would expect this effect to be biased towards horizon values of k_{\parallel} .

Power appears distributed in ‘pockets’ in the Stokes U power spectrum, not strongly correlated with the distribution of power in I. Stokes I is able to leak into Stokes U via D -term leakage (Thompson et al., 2017; Geil et al., 2011), which could occur at any post-amplification stage of observations, such as in cables or receivers. These leakages would be direction independent, and therefore uncorrelated in k -space. Such a mechanism could explain the behavior within Stokes U wedge. Before absolute calibration, similar structure is seen in the Stokes V power spectrum.

At these frequencies, Stokes V is thought to be intrinsically zero, with few exceptions. However, Hamaker et al. (1996) show that antennae rotated with respect to one another can produce erroneous Stokes V power via $I \rightarrow V$ leakage.³ This effect may explain some of the small pockets of power that remain in the Stokes V power spectrum after absolute calibration, although such an effect is also consistent with D -term leakage. The fact that power within the horizon was greater than the noise level may also have been due to $I \rightarrow V$ leakage through the primary beam.

The relationship between polarizations is highlighted in Figure 9.7. We show a slice of the wedges over $0.097 < k_{\perp} < 0.098 h\text{Mpc}^{-1}$ (~ 175 m) for Stokes I, Q, U and V (right panels) and the average power over these slices as a function of k_{\parallel} (left panel). The standard deviations for each Stokes parameter are shown as dotted lines. Dashed vertical lines show the horizon at $k_{\perp}=0.097$ (left) and super-horizon at $k_{\perp}=0.098$ (right).

A heartening aspect of Figure 9.7, and indeed all of the power spectra in this work, is that the power in Stokes Q, U and V proves to be just as confined within the horizon

³It should be noted that while such an error could plausibly have been made in the antenna placement for this imaging array, it is extremely unlikely that it would be made in the redundant PAPER configuration for EoR seasons. In these cases, the antennae were positioned to sub-cm accuracy.

as Stokes I. Whether the polarized Stokes parameters are due to real polarization or mis-calibration, not enough spectral structure is being introduced to move emission into the EoR window. Outside of the horizon, Stokes I, Q and U are consistent with the noise level expected for this range of k -modes ($P_{\text{noise}} \sim 10^9 \text{ mK}^2(h^{-1}\text{Mpc})^3$), according to the formalism Parsons et al. (2012a) and assuming a system temperature $T_{\text{sys}} = 450 \text{ K}$ (e.g. Moore et al., 2017).

9.3 Discussion and Conclusions

We have presented measurements of instrumental polarization leakage in PAPER-32 using 2D power spectra. These have allowed us to quantify some of the possible instrumental effects that could limit a statistical detection of the EoR within the wedge, diagnosed in the Fourier space most relevant EoR statistical detection experiments. To our knowledge, this is the first study of Q, U and V 2D power spectra at these k -modes. We have shown that power from Stokes Q, U and V is as confined to the wedge as Stokes I. Any calibration errors do not appear to spread power outside the horizon.

In their study of 2D power spectra, Asad et al. (2015) reported evidence of polarized leakage into the EoR window at the sub-percent level, considering a 4° degree field of view. Their study differs from this work not only over the field of view (4° versus almost whole-sky), but also in the observing mode (tracked versus drift scan) and in the different k -space probes by LOFAR’s longer baselines. In this work the power spectrum is calculated on a per-baseline basis, whereas their study calculates power spectra for gridded data which are more prone to mode mixing effects (Hazelton et al., 2013).

Our results are expected, in principle, to be more prone to leakage contamination due to the intrinsic extremely wide field of view of the PAPER primary beams, however, we see no evidence of leakage in the EoR window down to our sensitivity limits even without correction for polarized beams that is instead included in Asad et al. (2015). Our analysis indicates therefore that neither intrinsic polarized emission nor the PAPER primary

beam are leaking power in the EoR window, although longer integrations are required to demonstrate that this is true down to the sensitivities required for EoR detection.

We showed that systematics can be probed with high signal-to-noise using 2D polarized power spectra, using the inherently low polarization fraction at the frequencies PAPER observes at to our advantage. We found that gain errors on specific baselines were easily probed using Stokes Q power spectra. Gain errors appear as continuous streaks within the horizon at specific values of k_{\perp} , allowing us to diagnose the precision of the gain calibration on a per-baseline basis. This is much more difficult to do with only Stokes I power spectra in a non-redundant array, and can be accomplished quickly without imaging. While the features in the Stokes U power spectra are more difficult to attribute to specific baselines, they appear to be consistent with direction-independent leakage. Stokes V power is slightly higher than noise-level within the horizon, suggesting a small but unaccounted-for leakage term from Stokes I, an effect which was explored Nunhokee et al. (2017) – and found to be consistent with beam-leaked signal from Stokes I.




Figure 9.2: *Above:* Example of a Stokes I snapshot image (top left) with corresponding Stokes Q (top right), Stokes U (bottom left) and Stokes V (bottom right) images before absolute calibration. *Below:* The same organization as above, after absolute calibration. No primary beam correction was applied. The Stokes I image was deconvolved down to 5 Jy beam^{-1} whereas the other images were not deconvolved. Units are Jy beam^{-1} ; note the change in scale between polarizations and calibration stages.




Figure 9.3: The values of ionospheric RM for different lines of sight the range of LSTs in this analysis, as calculated by IONFR (Sotomayor-Beltran et al., 2013). The 21h, 0° line of sight goes beneath the horizon after LST=3h, and therefore has fewer data points.




Figure 9.4: The absolute value of delay-transformed visibilities over the bandwidth (146–166 MHz) used to create the power spectra shown in this Chapter. The left and right columns show the visibilities before and after absolute calibration (and for Stokes U and V, the application of the τ_{xy} parameter), respectively, for baseline formed by antennae 6 and 14 (~ 250 m in length, approximately East-West). The flux scale in the left column has been boosted for a more fair comparison to the absolute-calibrated data. From top to bottom, the rows correspond to Stokes I, Q, U and V. The horizon limit is marked by white dashed lines.




Figure 9.5: *Above:* Log-scaled 2D power spectra formed from (*top to bottom*): I, Q, U and V visibilities after before absolute calibration. Blank regions indicate the incomplete uv coverage for a given $k_{\perp}(u)$. The colorbar spans 10^9 to 10^{15} $\text{mK}^2(h^{-1}\text{Mpc})^3$. The flux scale has been boosted for a more fair comparison to the absolute-calibrated data. *Below:* The same organization as above, but after absolute calibration. Briefly, the structure in Stokes I is consistent with a point-source-dominated field with a weak diffuse component. The other Stokes parameters are consistent with calibration errors and systematics: Stokes Q shows gain errors on specific antennae, Stokes U gives an estimate of possible D-term leakage, and any structure in V shows unaccounted-for systematics, due to D-terms or mis-oriented antennae.




Figure 9.6: Log-scaled 2D power spectra formed from (*left to right*): I, Q, U and V absolute-calibrated visibilities. Blank regions indicate the incomplete *uv* coverage for a given $k_\perp(u)$. White and orange lines indicate the horizon and horizon plus a 50 ns boundary for super-horizon emission.




Figure 9.7: *Left:* The average power in $0.093 < k_{\perp} < 0.098 h\text{Mpc}^{-1}$ as a function of k_{\parallel} for each polarization. In the horizon–super-horizon range that we see the sharp fall-off in power indicative of the edge of the EoR window. Outside of the wedge, Stokes I, Q and U are at noise level, while Stokes V is below noise level, most likely due to the τ_{xy} calibration scheme removing a degree of freedom from this Stokes parameter. *Right:* the region of k-space that was averaged over to create the lines in the left panel. From top to bottom, the panels correspond to Stokes I, Q U and V.

Chapter 10

A view of the EoR window from the HERA-19 commissioning array

As emphasized in Chapter 9, it is important to constrain intrinsic and leaked polarized signal for any HI EoR experiment. The objective of this Chapter was an exploration of eight nights of data from the Hydrogen Epoch of Reionization Array (HERA) 19-element commissioning array, coupled with simulations of the instrument, in order to forecast how much of a problem polarization would pose for this interferometer. This work also represents the first power spectral analysis from HERA. While not in the realm of an EoR-level integration, we were able to offer some initial expectations for this new instrument’s performance in the Fourier domain.

This work is organized as follows: in Section 10.1 we review the theory behind polarization leakage into unpolarized signal and simulate the effect for a model of HERA. In Section 10.2 we describe the HERA data that we used, its calibration and reduction to power spectra. We present our results, and discuss the implications for HERA’s EoR measurements, in Section 10.3, and conclude in Section 10.4. We assume the cosmological parameters reported by Planck Collaboration et al. (2016) throughout.

10.1 Polarization Leakage Simulations

10.2 Observations & Reduction

In this work we used eight nights of observations from the HERA-19 commissioning array. HERA is a low-frequency interferometer composed of 14 m-diameter dishes arranged in a close-packed hexagonal array of 14.7 m spacing. The commissioning array consists of nineteen dishes (see Figure 10.1); HERA is being constructed in staged build-outs, and upon completion will consist of 350 dishes in a fractured hexagon configuration (see Dillon & Parsons, 2016; DeBoer et al., 2017). A feed cage containing two dipole feeds (recycled from the PAPER array, see Parsons et al. 2010), oriented in North-South and East-West directions, is suspended above each dish (Neben et al., 2016; Ewall-Wice et al., 2016a; Thyagarajan et al., 2016).

HERA only observes in drift-scan mode. The observations we used were eight nights, from Julian Date (JD) 2457548 to 2457555; LSTs 10.5 – 23 hr. Drift-scan visibilities were recorded every 10.7 seconds for 1024 evenly-spaced channels across the 100-200 MHz bandwidth. These data were divided into MIRIAD data sets roughly 10 minutes long. A night’s observation lasted 12 hours in total (6pm to 6am South African Standard Time; SAST); of these we used the central 10 hours, to avoid the thermal effects of the Sun.

To identify samples contaminated by radio frequency interference (RFI), a two-dimensional median filter in time and frequency was applied to the visibility data to smooth out high pixel-to-pixel variations, and remove significant outliers that were likely unphysical. The variance of the resulting data was computed, and points with a z -score greater than 6 (i.e., points where the value is more than 6σ away from the mean) were flagged as initial seeds for RFI extraction. A two-dimensional watershed algorithm was applied using these seeds as starting points, enlarging the regions of RFI-contamination to neighboring pixels with z -scores greater than 2, until all such pixels were flagged. Figure 10.2 shows the fractional RFI flag occupancy per time (displayed in LST) and frequency across the 8 days of observations. The majority of the band is relatively clear of RFI. Some clear features are: the




Figure 10.1: The perimeter of each dish in the HERA-19 array. A red "X" marks antennae that were identified during preprocessing and calibration as malfunctioning and were excluded from further analysis.

FM radio band (below 110 MHz), ORBCOMM satellite communications (137 MHz), an ISS downlink (150 MHz) and VHF TV channels (above 170 MHz)¹. The Galaxy, when transiting zenith at $LST \approx 17.75$ hours, is so bright that it appears to degrade our ability to flag RFI.

10.2.1 Calibration

HERA is designed to be calibrated using redundant calibration techniques (Dillon & Parsons, 2016), but for this preliminary view of HERA commissioning data, we used image-based calibration. Future studies with deeper integrations targeting EoR detections will take advantage of redundancy to obtain more precise calibration solutions (DeBoer et al., 2017). We used the CASA (McMullin et al., 2007) package for calibration, taking advantage of its CLEAN, gaincal and bandpass functions.

To enable the use of CASA, we first converted from native MIRIAD to a UVFITS

¹For an extended discussion of RFI as seen by HERA, see the public HERA Memo # 19




Figure 10.2: Fractional RFI flag occupancy per time and frequency over the eight days of observations.

file format which could be ingested by CASA using PYUVDATA (Hazelton et al., 2017). Using LSTs in which the Galactic center (GC; $\alpha, \delta = 17\text{h } 45\text{m } 40.04\text{s}, -29^\circ 0\text{m } 28.12\text{s}$) was transiting, we built a CLEAN model which modeled the GC as an unpolarized point source of strength 1 Jy and flat spectrum, which could be scaled appropriately later (see Equation 10.1). Clearly, this is an incomplete calibration model. However, as the objective of this work is to explore the response of the instrument in power spectrum space without combining baselines of different lengths, most of the purpose of the calibration is correcting an initial large cable delay per antenna. Treating the GC as unpolarized is adequate for this study. The large optical depth towards the GC (Oppermann et al., 2012) results in large amounts depolarization in the plane of the Galaxy (Wolleben et al., 2006). Moreover, we expected non-negligible amounts of beam depolarization due to the large solid angle of the synthesized beam.

For each night of observations, we used the CASA `gaincal` and `bandpass` functions to obtain frequency-dependent phase and amplitude solutions for each antenna and dipole arm. Four antennae had very deviant solutions, and their inclusion resulted in low-quality

images. These were omitted from further analysis (and are marked with red “X”s in Figure 10.1). Before calibration, we manually flagged the edges of the band (below 110 MHz and above 190 MHz), where spectral behavior is dominated by the high and low pass filtering in the HERA signal chain (DeBoer et al., 2017).

In Figure 10.3, we show images formed from the simulated pseudo-Stokes visibilities (top panels) and our observations (bottom panels). These are multi-frequency synthesis images, where we used all unflagged frequencies on either side of the band edges; 115 MHz to 188 MHz. We do not specify a beam model during imaging. At HERA’s position ((latitude, longitude) = (-30:43:17.5, 21:25:41.9)) the Galactic Center transits 2° from zenith, while the HERA primary beam has a FWHM of $\sim 5^{\circ}$ at 150 MHz (Neben et al., 2016). For the simulated visibilities, we flagged the same antennae as in the data. As expected for a compact array, the Stokes I images capture only a low-resolution view of the Galactic Center. The simulated and observed visibilities form remarkably similar images in Stokes I, Q and U, but the simulation under-predicts pseudo-Stokes V power. We defer further discussion to Section 10.3.

Example bandpass solutions from JD 2457548 are shown in Figure 10.4. Although some residual RFI remains obvious, the derived bandpasses were smooth. Thus, even though the gains were imprecise, we expected that using them should not add additional spectral structure.

The complex gain solutions were subsequently applied to the MIRIAD files. Figure 10.5 shows the effect of calibration on the visibilities of three nominally redundantly-spaced baselines. Shown in that figure are the phases of three V^{nm} visibilities from 14.7 m baselines before and after calibration. There were no shared antennae between the visibilities shown. The qualitative agreement is obvious, providing a consistency check on the solutions.

We did not attempt to calibrate D -terms in this work.

We down-selected to two relatively RFI-free 20 MHz sub-bands (Figure 10.2); 115 to 135 MHz and 152 to 172 MHz, henceforth referred to the “low band” and the “high





Figure 10.3: *Above:* Multi-frequency synthesis pseudo-Stokes images formed from simulation, where only a Stokes I sky was used; any polarized power is due to direction-dependent polarization leakage (see Section 10.1). *Below:* Multi-frequency synthesis pseudo-Stokes images formed from observed visibilities on JD 2457548. Both sets of panels show the Galactic Center (our calibrator source) close to transit in pseudo-Stokes I, Q, U and V visibilities (*top left, top right, lower left, lower right*). A Briggs-weighting with robustness 0 was used when gridding into the image plane. No deconvolution was performed. The colorbar is in units of Jy/Beam. A separate color scale is used for Stokes I for suitable dynamic range. An R.A., Dec. grid is shown, illustrating the wide-field nature of HERA observations.




Figure 10.4: Bandpass solutions for the North-South dipole orientation obtained for the functioning antennae in the array on JD 2457548. Differences in line color and style is merely to distinguish different antennae. Shaded regions indicate the effective sub-bands used for power spectrum analysis.

band". As we discuss in Section 10.2.2, these bands were multiplied by a Blackman-Harris window, centered on their central frequencies, before Fourier transforming in order to minimize side-lobes. This windowing lead to an noise-effective bandwidth of 10 MHz, appropriate for EoR analyses since the HI signal is to a reasonable approximation coeval over the corresponding redshift range (Furlanetto et al., 2006).

Pseudo-Stokes visibilities were formed from the instrumental polarizations. These visibilities were then scaled to the appropriate amplitude using a model for the GC spectrum

$$S_{\text{SgrA}^*}(v) \approx 3709 \text{ Jy} \times (v/408 \text{ MHz})^{-0.5} \quad (10.1)$$

drawn from the Global Sky Model (GSM; de Oliveira-Costa et al., 2008; Price, 2016a; Zheng et al., 2017). Note that the GSM is inherently $\sim 5\%$ uncertain at these frequencies. We note that this scaling is heavily resolution dependent; we are treating the Galactic Center as a point source when it is extended in reality. However, in Section 10.3 we show that we obtain sensible power levels for the foregrounds and noise, lending confidence to




Figure 10.5: The effect of calibration on the phases of visibilities from three redundantly-spaced 14.7 m baselines; nn polarization. The color scale is cyclic; black is $\pm\pi/2$ and white is 0 and $\pm\pi$. *Above:* before calibration; *below:* after calibration. A simple sky model was sufficient to enforce redundancy for redundant baselines.

our overall scaling.

10.2.2 Forming power spectra

Power spectra were formed according to the method used in Pober et al. (2013) and Kohn et al. (2016), which we briefly review here. All Fourier transforms were windowed using a Blackman-Harris window at the center of the sub-band, which minimized sidelobes. Parsons et al. (2012b) define the *delay transform* as the Fourier transform of a visibility for baseline ij and pseudo-Stokes parameter P along the frequency axis

$$\tilde{V}_{ij}^P(\tau, t) = \int dv \tilde{V}_{ij}^P(v, t) e^{2\pi i v \tau}. \quad (10.2)$$

We note that using a Blackman-Harris window will induce a correlation between consecutive τ modes. The Fourier transform of the window function in frequency will be sharply peaked in the delay space, and can be ignored to some extent. Hence the self-correlation of $V_{ij}^P(\tau, t)$ can be used to define the power spectrum, although the small correlation of different τ modes could effect the variance of the power spectrum (Parsons et al., 2014).

The power at each delay-mode and baseline can be represented in terms of their respective Fourier components k_{\parallel} and k_{\perp} (Parsons et al., 2012b; Thyagarajan et al., 2015a):

$$P(k_{\parallel}, k_{\perp}) \approx |\tilde{V}_{ij}^P(\tau)|^2 \frac{X^2 Y}{\Omega B} \left(\frac{c^2}{2k_B v^2} \right)^2,$$

$$k_{\parallel} = \frac{2\pi v_{21\text{cm}} H(z)}{c(1+z)^2} \tau,$$

$$k_{\perp} = \frac{2\pi}{D(z)\lambda} b$$

for: bandwidth B , angular area of the beam Ω , $v_{21\text{cm}} \approx 1420$ MHz, baseline length b ,

wavelength of observation λ , Hubble parameter $H(z)$, transverse comoving distance $D(z)$ and redshift-dependent scalars X and Y (Parsons et al., 2012a). Note that the angular area of the beam refers to the diagonal components of the Mueller matrices shown in Figure XXX. For further discussion of forming polarized power spectra in k -space, refer to Nunhokee et al. (2017).

To avoid a noise-bias when forming the $|\tilde{V}_{ij}^P(\tau, t)|^2$ term, we cross-multiplied consecutive integrations, rephasing the zenith angle of the latter to the former:

$$|\tilde{V}_{ij}^P(\tau, t)|^2 \approx |\tilde{V}_{ij}^P(\tau, t) \times \tilde{V}_{ij}^P(\tau, t + \Delta t) e^{i\theta_{ij,\text{zen}}(\Delta t)}| \quad (10.3)$$

where $\theta_{ij,\text{zen}}(\Delta t)$ was the appropriate phasing for baseline ij and $\Delta t = 10.7$ seconds.

Pseudo-stokes power spectra were formed for each pair of integrations, for every baseline. After forming power spectra, baselines of identical lengths were averaged together. Appealing to cosmological isotropy, baselines of the same length but different orientation should be sampling the same cosmological structure. These 2D power spectra were averaged over our 8 days of observations. Note that all averaging was performed after forming power spectra; this incoherent averaging was non-optimal from a signal-to-noise perspective outside the wedge, slightly reducing our sensitivity in the EoR window. However, the intention of this investigation was not a deep integration on noise; we were more interested in the polarized response of the instrument. As such, the power spectra presented in the Section below should be interpreted as approximate.

10.3 Results & Discussion

Power spectra are shown for the high and low bands in Figure 10.6 and Figure 10.7, respectively, where white dotted lines mark the boundary of the EoR window on the 2D plots. The same data are presented in middle and lower panels, with the latter overlaid as lines to emphasize common features of the power spectra with respect to baseline length.

Theoretical noise levels for the high and low bands were between $P_{\text{noise}}(k) \approx 1.7 \times 10^8$

$\text{mK}^2\text{Mpc}^3\text{h}^{-3}$ and $3.4 \times 10^9 \text{ mK}^2\text{Mpc}^3\text{h}^{-3}$ in the high band, and between $2.3 \times 10^8 \text{ mK}^2\text{Mpc}^3\text{h}^{-3}$ and $6.1 \times 10^9 \text{ mK}^2\text{Mpc}^3\text{h}^{-3}$ in the low band. These estimates used using a temperature model of the sky

$$T_{\text{sky}} = 180\text{K} \left(\frac{\nu}{180\text{MHz}} \right)^{-2.55}, \quad (10.4)$$

assume receiver temperatures of 300 K and 600 K for the high band and low band, respectively (DeBoer et al., 2017, also see the public HERA Memo # 16), and were calculated according to the formalism for noise power spectra in Parsons et al. (2012b), with the inclusion of a baseline-number dependence (to account for different occupancies in each k_{\perp} bin). These noise power estimates were roughly corroborated by our observations (see Figure 10.8). We observe excess noise on the shortest baselines (also obvious in the lower panels of Figures 10.6 and 10.7).

10.3.1 General features of the power spectra

The most striking feature of these power spectra is the degree of foreground isolation achieved in all pseudo-Stokes parameters. In similar studies of 2D polarized power spectra, both PAPER (Kohn et al., 2016) and LOFAR (Asad et al., 2017) measurements found “filled” regions of Fourier space out to the edge of the EoR window (in the delay-spectrum paradigm, this corresponds to the horizon; zenith angle $\pm 90^\circ$), with some supra-horizon leakage (Pober et al., 2013) into the EoR window itself. The power spectra in Figures 10.6 and 10.7 show no such behavior; all foreground emission appears to be contained within a narrow region around $k_{\parallel} = 0 \text{ h/Mpc}$. This behavior was predicted for an array of HERA-like dishes by Thyagarajan et al. 2015a (although that study only concentrated on the Stokes I component).

Power at horizon delays, as predicted by Thyagarajan et al. (2015a) and Neben et al. (2016), was not observed. This was likely a resolution effect. To resolve horizon-delay power, one would need to sample many periods of $\tau_h = b/c$, where b is the magnitude of the baseline vector. The maximum length baseline in the HERA-19 array was 58.4 m, cor-

responding to a ~ 5 MHz period: barely sampled by the 10 MHz windows we use in this study. The lack of horizon power is corroborated by the simulations of the HERA delay response in Ewall-Wice et al. (2016a) and Thyagarajan et al. (2016), although those studies used a different windowing function for the delay transform. Their simulations also predict a high degree of foreground isolation: the presence of noise in our data of course meant that we do not realize the 11 dex of isolation that can be achieved in simulation, but the ~ 8 dex we do see, without any foreground subtraction and a simple calibration, speaks to the power of HERA’s future capabilities.

Visible in the observational data, but not in the simulation, is an excess of power at $k_{\parallel} = 0.04$ h/Mpc, corresponding to a delay of 100 ns, which is independent of baseline length. This suggests that its origins are in the HERA signal chain. There are 15 m coaxial cables at one stage of the signal chain from the HERA dishes to the correlator². In the limit of little delay induced by the cable and our limited delay resolution, a reflection along this stage of the signal chain would produce an alias of the foreground signal at a $\tau \approx 100$ ns (Beardsley et al., 2016; Ewall-Wice et al., 2016b).

10.3.2 Day-to-day variability

The foreground and EoR window power levels appeared to be relatively stable between days, with variation most likely due to the incomplete sky model used for gain calibration. Figure 10.8 shows power as a function of baseline length for $k_{\parallel} = 0$ h/Mpc (solid lines) and $k_{\parallel} = 0.2$ h/Mpc (dot-dashed lines). Deviations from the mean at $k_{\parallel} = 0$ h/Mpc may be a limitation imposed by our simplistic sky model. Since the noise levels in the EoR window region remained noise-like throughout our observations, the uncertainty in the absolute gain scale did not have a large impact on our largely-diagnostic investigation.

²This stage of the signal chain is only present in the commissioning array. Future HERA build-outs will transition to a different architecture (DeBoer et al., 2017).

10.3.3 Polarimetric results

Figures 10.6 and 10.7 qualitatively illustrate that the simulations described in Section 10.1 reproduced the main features of the observed power spectra. The simulations were run only with a Stokes I sky component and no simulated calibration errors, so the only signal in the polarized power spectra was from wide-field beam leakage (Figure XXX). An example comparison between simulation and observation in the image plane is shown in Figure 10.3.

In Figure 10.9 we show the power levels observed on the shortest baseline (14.7 m) compared to our simulations for each band. The simulations used an unpolarized diffuse sky model (the most recent version of the GSM; Zheng et al., 2017), which should be accurate at the scales probed by a 14.7 m baseline. Inset panels zoom-in on the region around $k_{\parallel} = 0 \text{ h/Mpc}$, where most of the foreground power was concentrated. We saw that the simulations reproduced $\sim 75\%$ of the foreground power observed in pseudo-Stokes I in the high band, and over-predicted foreground power by $\sim 35\%$ in the low band. This could have been due to unrealistic frequency scaling of the diffuse foregrounds in the GSM.

For pseudo-Stokes Q and U, the simulations accounted for $\sim 60 - 75\%$ of power seen within the pitchfork region, suggesting that most of the power seen in these power spectra, at least for the shortest baselines, can be mostly attributed to direction-dependent leakage effects. As noted in Section 10.1, residual gain and phase errors are able to leak a fraction of pseudo-Stokes I into Q and U, but some fraction of the observed power ($\leq 25\%$) may have been due to linearly polarized foregrounds. This is corroborated by residual power close to the location of the Galactic Center, and increased power over the sky, in the observed pseudo-Stokes Q and U skies compared to the simulated ones in Figure 10.3. As the Galactic Center is the highest-amplitude source of power, we expect residual gain errors to be most obvious in the same position as it is in the pseudo-Stokes I image. Such an excess is present in the observed pseudo-Stokes Q and U images, but absent in the simulated ones – pointing to direction-independent gain errors being present.

However, the simulated pseudo-Stokes Q and U images contain only direction-dependent leakage from Stokes I. Since they reproduce most of the features seen in the observed data, pseudo-Stokes Q and U are clearly dominated by direction dependent leakage.

Lenc et al. (2016) observed linearly polarized emission from diffuse structure with $\sim 1.6 - 4.5\%$ fractional polarization at 150 MHz, corresponding to power levels of $\sim 10^5 \text{ mK}^2 \text{Mpc}^3 \text{h}^{-3}$. This power level is similar to expected EoR power levels (e.g. Lidz et al., 2007; Moore et al., 2013; Nunhokee et al., 2017); a detection of a power spectrum of polarized galactic synchrotron will require much deeper integrations.

The observed pseudo-Stokes V power spectrum was more poorly modelled by our simulation. In both bands we observed ~ 20 dB more power in pseudo-Stokes V at $k_{\parallel} = 0 \text{ h/Mpc}$ than predicted by our simulations. The peak power observed in pseudo-Stokes V was roughly 0.1% of the peak power observed in pseudo-Stokes I. Likewise in the sky images shown in Figure 10.3, there is little pseudo-Stokes V power in the simulated images, compared to observation. This suggests that most or all of the power in pseudo-Stokes V is due to direction independent leakage. While the leakage appears localized in Figure 10.3, we see in Figure 10.9 that it is statistically similar to pseudo-Stokes Q and U in power. Since D -terms cause direction-independent leakage from pseudo-Stokes I to pseudo-Stokes V, the excess power we observed could be interpreted as an approximate D -term level of $\sim 1\%$ (Thompson et al., 2017). This is similar to D -term levels from other low frequency instruments such as MWA-32, which was found to have $\sim 2\%$ D -terms (G. Bernardi, private communication). The under-prediction of pseudo-Stokes V from the simulation could, of course, also be due to some unmodelled direction-dependent instrumental effect.

To understand which effect, if either, is dominant, a precise D -term calibration of HERA is required. This effort is underway with data taken with bright polarized point sources in transit, and will be presented in future work. Another potential cause of the discrepancy could have been that our simulations under-predicted Stokes V power, due to lack of accounting for some variety of instrumental circular polarization.

In Section 10.3.1 we noted the presence of excess power at $k_{\parallel} = \pm 0.04 \text{ h/Mpc}$ ($\pm 100 \text{ ns}$) that was independent of baseline length, suggesting that it was due to a reflection along 15 m cables. Figure 10.9 shows that power at this delay is not consistent between polarizations. Stokes U and V power only exhibited excess signal at -100 ns in the high band, and in the low band, it was only Stokes U that did not exhibit that excess at +100 ns. This may be a clue about the polarization state of cable reflections, perhaps as a function of frequency, but we defer this to future work – noting it as a point of interest here.

10.4 Conclusions

In this work we have presented polarized power spectra from the HERA-19 commissioning array. With modest calibration, HERA is able isolate total intensity and polarized foregrounds to within the “pitchfork” region of k -space, as predicted by Thyagarajan et al. (2015a), lending confidence to its future performance as an instrument capable of both detecting and characterizing the EoR power spectrum. Of course, the array used in this study had just 19 antennae, 15 of which were used for analysis – future build-outs of HERA with up to 350 antennae will require strong quality-assurance efforts.

Simulations of the polarized response of the instrument, mapped into the same Fourier space as the data, suggest that most or all of the polarized power observed in pseudo-Stokes Q and U power spectra is due to direction-dependent beam leakage from pseudo-Stokes I. Residual gain and phase errors could account for the rest of the power, but some fraction of the total ($\leq 25\%$) may be due to linearly polarized foregrounds. Excess power in pseudo-Stokes V may be due to D -terms at the 1% level, but a full image-based calibration with a polarized point source is required to confirm this. The general accuracy of our simulations suggests current modelling of the complex HERA beam is accurate.




Figure 10.6: Results from the high-band (157–167 MHz). White dotted lines indicate the boundary of the pitchfork and the EoR window. A black dotted line indicates the $k_{\parallel} = 0$ h/Mpc line. *Top:* Simulated power spectra in Stokes I, Q, U and V, following the formalism in Section 10.1 – no polarized sky model was used, so power in Stokes Q, U and V was only due to direction-dependent leakage from Stokes I. No instrumental noise was included in the simulation. *Middle:* Eight-day average power spectra from data. *Bottom:* The same data as shown in the middle panel, but with each baseline length overlaid on one another to allow shared features to be more easily identified.




Figure 10.7: Results from the low-band (120–130 MHz), arranged in the same format as Figure 10.6.




Figure 10.8: High band power as a function of baseline length for the center of the pitch-fork ($k_{\parallel} = 0 \text{ h/Mpc}$; solid lines) and in the EoR window ($k_{\parallel} = 0.2 \text{ h/Mpc}$; dot-dashed lines) for each JD of observation. The black dashed line represents the approximate noise power assuming a receiver temperature of 300 K. A very similar relationship is shown in the low band, but with a higher noise floor, which is consistent with system temperature as a function of frequency. The noise level climbs with baseline length as the compact nature of the array gives more short baselines to average-over in a given $(k_{\parallel}, k_{\perp})$ bin than longer ones.




Figure 10.9: Simulated and observed power as a function of k_{\parallel} for the shortest baseline (14.7 m). *Right to left*: pseudo-Stokes I, Q, U and V; *above*: the high band; *below*: the low band. The simulations were noiseless and used an unpolarized sky model. Inset panels zoom-in on the peak region. They capture the foreground power levels in pseudo-Stokes I, Q and U, suggesting all power in Q and U is due to leakage from Stokes I. The power level in V is highly discrepant, however, suggesting some sort of beam-independent instrumental leakage.

Chapter 11

Deep integrations on polarization with PAPER-128

In Chapter 9, I presented polarized power spectra from a short integration – a few hours of one night – over a wide range of k_{\perp} -modes probed by the PAPER-32 polarized imaging array (Kohn et al., 2016). Chapter 10 presented the first power spectral results from HERA. The HERA-19 commissioning array was small and dense, meaning that only a few k_{\perp} -modes were accessible. For that study, we averaged over 10 hours per night, for 8 consecutive nights (Kohn et al. 2018). This Chapter presents results from the PAPER-128 array. In this Chapter I present roughly one quarter of the total number of observations recorded by this interferometer (Section 11.1). I show results of a deep integration on a very narrow range of k_{\perp} -modes (corresponding to ~ 30 m spacings of the redundant grid; Section 11.2) and discuss the implications for deep, fully-polarized integrations with large interferometers (Section 11.3).

11.1 Observations & Reduction

PAPER-128 was the largest build-out of the PAPER experiment. As described in Chapter 4, PAPER-128 consisted of 128 antennas, 112 of which were arranged in a redundant

Table 11.1. PAPER-128 Observing Seasons & Epochs

Season	Epoch	Julian Dates	Calendar Dates	Notes
1	1	2456617 - 2456673	Nov 20, 2013 - Jan 15, 2014	1/8 F-Engine failure
	2	2456678 - 2456724	Jan 20, 2014 - Mar 7, 2014	
2	1	2456625 - 2456732	Mar 8, 2014 - Mar 7, 2014	Too few data
	2	2456843 - 2456873	Jul 4, 2014 - Aug 3, 2014	Uninteresting LST range
	3	2456881 - 2456928	Aug 11, 2014 - Sep 27, 2014	Many malfunctioning antennas
	4	2456942 - 2457008	Oct 11, 2014 - Dec 11, 2014	
	5	2457030 - 2457050	Jan 7, 2015 - Jan 27, 2015	Many malfunctioning antennas

grid. An annotated photograph of the array is shown in Figure 11.1. In this section, I review the PAPER-128 campaign (Section 11.1.1) and the subsequent reduction of roughly one quarter of the total number of observations (Section 11.1.2).

11.1.1 Overview of PAPER-128 observations

Observations were recorded for two years, with first light on November 20th 2013 and final readings on January 27th 2015. However, these observations were not always contiguous. Human errors, experimentation and malfunctioning electronics required the correlator and connected electronics to be turned off and restarted, altering the characteristic phasing and gain scale of the array. Each of these restarts constituted the beginning of a new “Epoch” of the array which required different quality assurance steps and initial calibration stages. Table 11.1.1 summarizes the length and nature of these Epochs.

Figure 11.2 illustrates the challenge imposed by the correlator restarts. Epoch changes were characterized by large shifts in the overall phase of visibilities, of course leading to changed magnitudes of the real and imaginary parts of the visibilities recorded. The quality assurance metrics described in Chapter 6 were sensitive to these changes. Of course, absolute calibration – that is phasing to the correct point on the sky and scaling the visibilities from arbitrary to physical flux density units – had to be run separately on




Figure 11.1: An annotated photograph of the PAPER-128 array, looking to the East. Highlighted are the 112 antenna redundant grid, with 15 m East-West spacings between each row; outlier antennas from the main grid used to increase *uv*-coverage; coaxial cables running to the receiverators and correlator (see Chapter 4). An inset panel shows a PAPER sleeved dipole. Photo credit: J. E. Aguirre. Figure credit: C. D. Nunhokee; (Nunhokee, 2018).

individual Epochs.

As shown in Table 11.1.1, Season 2 had a larger number of observed nights than Season 1. However, the analysis of Season 2 was especially challenging due to large numbers of malfunctioning antennas. This may have been due to the antennas ageing past a critical point. Nothing in the array was replaced during build-outs except for the correlator; 32 of the antennas had been out in the desert for 4–5 years (PAPER-32) and another 32 for 3–4 years (PAPER-64). Most of the time, these were the antennas that were identified as malfunctioning.

Season 2 Epoch 4 was relatively well-behaved, and may be analyzed in the future. For this work, we concentrated our analysis on Season 1. Season 1 Epoch 2 was ten days shorter than Season 1 Epoch 1, but Epoch 1 had two major challenges associated with it: a data loss event, and an F-engine failure. Due to human error, Epoch 1 data was deleted and had to be restored and recompressed (see Chapter 6). This was almost entirely successful, at the loss of one week’s worth of observations. The F-engine failure was more critical. We discovered during our analysis that exactly one eighth of the antennas in the array produced noise-like visibilities with the rest of the array, but normal correlations between one another. These antennas had “seceded” from the array. They shared the characteristic of all being attached to the same F-engine (of which there were eight; see Figure 4.2). This suggested that there was a clock-offset on that F-engine, resulting in no correlation between the signal from those antennas and those running through the other in-sync F-engines.

This left Season 1 Epoch 2 as the most well-characterized and well-behaved Epoch of PAPER-128 observations, and we focused on this Epoch alone from now on. Cross-polarization metrics identified seven incorrectly-rotated antennas, which were corrected during initial processing. Eight antennas were identified as malfunctioning by mean visibility amplitude metrics, including one of those that was incorrectly-rotated. The state of the array for Season 1 Epoch 2 is summarized in Figure 11.3.




Figure 11.2: The challenge of Epoch changes. Shown are all of the Season 2 time samples of the visibilities recorded by the 30 m baseline between antennas 1 and 4, LSTs 0–5, for only the 150 MHz frequency bin. The above panel shows a histogram of the real part of the visibilities as a function of LST – there is a dramatic change in magnitude with respect to LST. Likewise, the lower panel shows the phase of each visibility sample (color axis) as a function of LST (vertical axis) and day of observation (x axis). There are obvious large shifts in phase, which require separate calibration stages.




Figure 11.3: Good, bad (red crosses) and rotated (blue circles) antennas in Season 1 Epoch 2, as identified by the cross-polarization and visibility amplitude metrics defined in Chapter 6.

11.1.2 Reduction of Season 1 Epoch 2 data

After correcting for cross-polarized and malfunctioning antennas, we

11.2 Results

11.3 Discussion & Conclusions

Part III

Expanding the potential of EoR

measurements

Chapter 12

Time-Averaged Visibilities

Chapter 13

Higher-order correlation functions between kSZ and 21cm observations

Chapter 14

Deep Learning for 21cm Observations

Chapter 15

Conclusions

Appendices

Appendix A

Software

Software engineering and maintenance of existing codebases has been, generally speaking, historically undervalued and unappreciated by the astronomy community (Muna et al., 2016). In this Appendix I would like to provide a brief description of the major software packages used in this work – without which, the work would not exist.

A.1 Astronomical Interferometry in Python (aipy)

The aipy software package (Parsons, 2016) was developed by a team based largely at the University of California, Berkeley and led by Aaron Parsons. Developed under NSF funding for the PAPER experiment, it provides a Python API to interact with interferometric visibilities stored in the MIRIAD file format (Sault et al., 2011). It is able to efficiently query large MIRIAD files due the APIs closeness to the underlying C code. It also contains calibration, deconvolution, imaging and phasing code in Python, and interfaces with HEALPix (see Section A.5, below) as well as other astronomical Python packages.

aipy is maintained by the HERA software team, and can be found at: <https://github.com/HERA-Team/aipy>.

A.2 Astronomy in Python (astropy)

astropy is an open-source and community-developed core Python package for Astronomy, containing a host of extremely useful utility functions and objects (Astropy Collaboration et al., 2013).

A.3 Common Astronomy Software Applications (CASA)

CASA is under active development, with the primary goal of supporting the data post-processing needs of the next generation of radio telescopes. It is developed by an international consortium of scientists based at the National Radio Astronomical Observatory (NRAO), the European Southern Observatory (ESO), the National Astronomical Observatory of Japan (NAOJ), the CSIRO Australia Telescope National Facility (CSIRO/ATNF), and the Netherlands Institute for Radio Astronomy (ASTRON), under the guidance of NRAO (McMullin et al., 2007).

A.4 Deep Learning packages

Experimentation with deep learning analyses of 21 cm simulated observations took place in Keras (Chollet et al., 2015), PyTorch (Paszke et al., 2017) and Tensorflow (Abadi et al., 2016).

A.5 Hierarchical Equal Area isoLatitude Pixelization of the sphere (HEALPix)

The HEALPix software, and its Python wrapper healpy, provide a pixelization which subdivides a spherical surface into pixels which each cover the same surface area as every other pixel. Pixel centers occur on a discrete number of rings of constant latitude. This

scheme makes natively spherical measurements, such as angular power spectra and wide-field images, simple and efficient to interact with (Górski et al., 2005).

A.6 pyuvdata

pyuvdata provides a Python interface to interferometric data. It can read and write MIRIAD and UVFITS file formats, as well as read CASA measurement sets and FHD (Sullivan et al., 2012b) visibility save files (Hazelton et al., 2017).

pyuvdata is maintained by the HERA software team, and can be found at: <https://github.com/HERA-Team/pyuvdata>.

A.7 The Scientific Python Ecosystem (scipy)

Many of the above tools require at least one of the many packages under the `scipy` ecosystem. It is truly foundational to almost any scientific analysis that takes place in Python (Jones et al., 2001).

Bibliography

Abadi, M. et al. 2016, ArXiv e-prints

Afraimovich, E. L., Astafyeva, E. I., Oinats, A. V., Yasukevich, Y. V., & Zhivetiev, I. V. 2008, *Annales Geophysicae*, 26, 335

Ali, Z. S. et al. 2015, ArXiv e-prints

Appleton, E. V. 1946, *Nature*, 157, 691

Asad, K. M. B., Koopmans, L. V. E., Jelić, V., de Bruyn, A. G., Pandey, V. N., & Gehlot, B. K. 2017, ArXiv e-prints

Asad, K. M. B. et al. 2016, *MNRAS*, 462, 4482

—. 2015, *MNRAS*, 451, 3709

Astropy Collaboration et al. 2013, *A& A*, 558, A33

Backer, D. C. et al. 2010, HERA âš Hydrogen Epoch of Reionization Arrays, Submitted for consideration by the Astro2010 Decadal Survey Program Panel

Beardsley, A. P. et al. 2016, *ApJ*, 833, 102

Bernardi, G. et al. 2013, *ApJ*, 771, 105

Brunner, F. K. & Welsch, W. M. 1993, *GPS World*, 4, 42

Carilli, C. L. & Rawlings, S. 2004, *NewAR*, 48, 979

Carozzi, T. D. & Woan, G. 2009, MNRAS, 395

Chollet, F. et al. 2015, Keras, <https://github.com/fchollet/keras>

Civil Aviation Authority. 2016, <http://wwwcaa.co.za/Pages/Search-results.aspx?k=frequency#k=113.3%20MHz>, website

Condon, J. J., Cotton, W. D., Greisen, E. W., Yin, Q. F., Perley, R. A., Taylor, G. B., & Broderick, J. J. 1998, AJ, 115, 1693

Cornwell, T. 2016

Cotton, W. & Condon, J. 2002, in Proceedings of the URSI General Assembly

de Oliveira-Costa, A., Tegmark, M., Gaensler, B. M., Jonas, J., Landecker, T. L., & Reich, P. 2008, MNRAS, 388, 247

DeBoer, D. R. et al. 2017, PASP, 129, 045001

Dewdney, P. E., Hall, P. J., Schilizzi, R. T., & Lazio, T. J. L. 2009, Proceedings of the IEEE, 97, 1482

Dillon, J. S. et al. 2017, ArXiv e-prints

—. 2014, Phys. Rev. D, 89, 023002

—. 2015, Phys. Rev. D, 91, 123011

Dillon, J. S. & Parsons, A. R. 2016, ApJ, 826, 181

Erdogan, E. et al. 2016, in EGU General Assembly Conference Abstracts, Vol. 18, EGU General Assembly Conference Abstracts, 12685

Erickson, W. C., Perley, R. A., Flatters, C., & Kassim, N. E. 2001, A&A, 366, 1071

Ewall-Wice, A. et al. 2016a, ApJ, 831, 196

- . 2016b, MNRAS, 460, 4320
- Finlay, C. C. et al. 2010, Geophysical Journal International, 183, 1216
- Furlanetto, S. R., Oh, S. P., & Briggs, F. H. 2006, Physics Review, 433, 181
- Geil, P. M., Gaensler, B. M., & Wyithe, J. S. B. 2011, MNRAS, 418, 516
- Górski, K. M., Hivon, E., Banday, A. J., Wandelt, B. D., Hansen, F. K., Reinecke, M., & Bartelmann, M. 2005, ApJ, 622, 759
- Hamaker, J. P., Bregman, J. D., & Sault, R. J. 1996, A&AS, 117, 137
- Hazelton, B., Beardsley, A., Pober, J., Jacobs, D., Ali, Z., & Lanman, A. 2017, HERA-Team/pyuvdata: Version 1.1
- Hazelton, B. J., Morales, M. F., & Sullivan, I. S. 2013, ApJ, 770, 156
- Hurley-Walker, N. et al. 2017, MNRAS, 464, 1146
- Iglewicz, B. & Hoaglin, D. 2010, Engineering Statistical Handbook, 1
- Iijima, B. et al. 1999, Journal of Atmospheric and Solar-Terrestrial Physics, 61, 1205
- Intema, H., Van der Tol, S., Cotton, W., Cohen, A., Van Bemmel, I., & Röttgering, H. 2009, Astronomy & Astrophysics, 501, 1185
- Jacobs, D. C. et al. 2016, ApJ, 825, 114
- . 2013, ApJ, 776, 108
- . 2015, ApJ, 801, 51
- Jelić, V. et al. 2014, A&A, 568, A101
- . 2015, A&A, 583, A137

- Jones, E., Oliphant, T., Peterson, P., et al. 2001, SciPy: Open source scientific tools for Python, [Online; accessed <today>]
- Klobuchar, J. A. 1983, Ionospheric Effects on Earth-Space Propagation
- Kohn, S. A. et al. 2016, ApJ, 823, 88
- Komjathy, A., Sparks, L., Wilson, B. D., & Mannucci, A. J. 2005, Radio Science, 40
- Lanyi, G. E. & Roth, T. 1988, Radio Science, 23, 483
- Lenc, E. et al. 2017, PA&A, 34, e040
- . 2016, ApJ, 830, 38
- Lidz, A., Zahn, O., McQuinn, M., Zaldarriaga, M., Dutta, S., & Hernquist, L. 2007, ApJ, 659, 865
- Liu, A., Tegmark, M., Morrison, S., Lutomirski, A., & Zaldarriaga, M. 2010, MNRAS, 408, 1029
- Loi, S. T. et al. 2015, Geophys. Res. Lett., 42, 3707
- MacDoran, P. & Spitzmesser, D. 1989, Method and apparatus for deriving pseudo range from earth-orbiting satellites, uS Patent 4,797,677
- Mannucci, A., Wilson, B., Yuan, D., Ho, C., Lindqwister, U., & Runge, T. 1998, Radio science, 33, 565
- McMullin, J. P., Waters, B., Schiebel, D., Young, W., & Golap, K. 2007, in Astronomical Data Analysis Software and Systems XVI, ed. R. A. Shaw, F. Hill, & D. J. Bell, 376, ASP, San Francisco, CA, 127
- Mevius, M. et al. 2016, Radio Science, 51, 927
- Mohan, N. & Rafferty, D. 2015, PyBDSM: Python Blob Detection and Source Measurement, Astrophysics Source Code Library, record 1502.007

- Moore, D. F. 2014, PhD thesis, University of Pennsylvania
- Moore, D. F. et al. 2017, ApJ, 836, 154
- Moore, D. F., Aguirre, J. E., Parsons, A. R., Jacobs, D. C., & Pober, J. C. 2013, ApJ, 769, 154
- Morales, M. F., Hazelton, B., Sullivan, I., & Beardsley, A. 2012, ApJ, 752, 137
- Muna, D. et al. 2016, ArXiv e-prints
- Neben, A. R. et al. 2016, ApJ, 826, 199
- NRC, N. R. C. 2010, New Worlds, New Horizons in Astronomy and Astrophysics (Washington, DC: The National Academies Press)
- Nunhokee, C. D. 2018, PhD thesis, Rhodes University
- Nunhokee, C. D. et al. 2017, ApJ, 848, 47
- Offringa, A. R., van de Gronde, J. J., & Roerdink, J. B. T. M. 2012, A&A, 539, A95
- Oppermann, N. et al. 2012, A& A, 542, A93
- Parsons, A. 2016, AIPY: Astronomical Interferometry in PYthon, Astrophysics Source Code Library
- Parsons, A. et al. 2008, PASA, 120, 1207
- Parsons, A., Pober, J., McQuinn, M., Jacobs, D., & Aguirre, J. 2012a, ApJ, 753, 81
- Parsons, A. R. & Backer, D. C. 2009, AJ, 138, 219
- Parsons, A. R. et al. 2010, AJ, 139, 1468
- . 2014, ApJ, 788, 106

- Parsons, A. R., Pober, J. C., Aguirre, J. E., Carilli, C. L., Jacobs, D. C., & Moore, D. F. 2012b, ApJ, 756, 165
- Paszke, A. et al. 2017
- Patil, A. H. et al. 2017, ApJ, 838, 65
- Pen, U.-L., Chang, T.-C., Hirata, C. M., Peterson, J. B., Roy, J., Gupta, Y., Odegova, J., & Sigurdson, K. 2009, MNRAS, 399, 181
- Planck Collaboration et al. 2016, A& A, 594, A13
- Pober, J. 2015, Updated 21cm Experiment Sensitivities, Tech. Rep. 4, Brown University
- Pober, J. C. et al. 2015, ApJ, 809, 62
- . 2013, ApJL, 768, L36
- . 2012, AJ, 143, 53
- Price, D. C. 2016a, PyGSM: Python interface to the Global Sky Model, Astrophysics Source Code Library
- . 2016b, ArXiv e-prints
- Royden, H., Miller, R., & Buennagel, L. 1984, Radio science, 19, 798
- Sault, R. J., Hamaker, J. P., & Bregman, J. D. 1996, A&AS, 117, 149
- Sault, R. J., Teuben, P., & Wright, M. C. H. 2011, MIRIAD: Multi-channel Image Reconstruction, Image Analysis, and Display, Astrophysics Source Code Library
- Schilizzi, R. et al. 2007, SKA Memorandum, 100
- SKA Panel. 2012, Expert Panel on Radio Quiet Zone and RFI Regulation Report, Tech. rep., Square Kilometer Array, <https://www.skatelescope.org/wp->

content/uploads/2012/06/79_Report_of_The_Expert_Panel_on_Radio_Quiet_Zone_and_RFI_Regulation.pdf

Smirnov, O. M. 2011, A&A, 527, A106

Sotomayor-Beltran, C. et al. 2013, A&A, 552, A58

South African Civil Aviation Authority. 2008, FLIGHT CALIBRATION DATES OF VOR's IN THE REPUBLIC OF SOUTH AFRICA, Tech. Rep. 25.6, Aeronautical Information Circular, <http://209.203.9.244/resource%20center/AIC'S/25.6.pdf>

Staatskoerant. 2008, South African Table of Frequency Allocations, Tech. Rep. 31264, ICASA, http://thornton.co.za/resources/31264_890_complete-1.pdf

Stefan, I. I. et al. 2013, MNRAS, 432, 1285

Sullivan, I. S. et al. 2012a, ApJ, 759, 17

—. 2012b, ApJ, 759, 17

Tariku, Y. A. 2015, Earth, Planets, and Space, 67, 35

Thompson, A. R., Moran, J. M., & Swenson, Jr., G. W. 2017, Interferometry and Synthesis in Radio Astronomy, 3rd Edition

Thyagarajan, N. et al. 2015a, ApJ, 804, 14

—. 2015b, ApJ, 807, L28

Thyagarajan, N., Parsons, A. R., DeBoer, D. R., Bowman, J. D., Ewall-Wice, A. M., Neben, A. R., & Patra, N. 2016, ApJ, 825, 9

Tingay, S. J. et al. 2013, PASA, 30, 7

Titheridge, J. E. 1972, P& SS, 20, 353

Trott, C. M. et al. 2016, ApJ, 818, 139

van Haarlem, M. P. et al. 2013, A&A, 556, A2

Vedantham, H. K. & Koopmans, L. V. E. 2015, MNRAS, 453, 925

—. 2016, MNRAS, 458, 3099

Wieringa, M. H. 1992, Experimental Astronomy, 2, 203

Wolleben, M., Landecker, T. L., Reich, W., & Wielebinski, R. 2006, A& A, 448, 411

World Aero Data. 2016, http://worldaerodata.com/wad.cgi?nav=VICTORIA+WEST&nav_id=VWW&nav_id=VWW&nav_id=VWW
website

Yatawatta, S. et al. 2013, A&A, 550, A136

Zheng, H. et al. 2014, 445, 1084

Zheng, H. et al. 2017, MNRAS, 464, 3486

Zolesi, B. & Cander, L. R. 2014, Ionospheric prediction and forecasting (Springer)